

Ally Quinn Story Backbone

Consensus Brief for Claude Review

Prepared for independent model-to-model adjudication

24 August 2026

Contents

Decision requested	1
Provisional answer	2
Authority and limits	2
Evidence that governs the decision	2
Where the two positions already agree	2
Claims that need amendment	3
1. "A closed case is public"	3
2. "The antagonist does not need to act"	3
3. "The clock problem ends"	4
4. "A killer walking around Havenbay creates permanent jeopardy"	4
Proposed shared backbone	4
Canon wording candidate	4
Four-part episode engine	4
Fresh cases remain possible	5
Two-axis case taxonomy	5
Axis one: What is wrong with the record?	5
Axis two: What is happening now?	5
Proposed jeopardy rule for adjudication	5
Candidate replacement	6
Cheapest useful test	6
Recommended falsification test	6
Decision rule without scores	6
Implications for the current slate	6
Document impact if adopted	7
Questions Claude must adjudicate	7
Consensus protocol	7
Revision history	8

Decision requested

Should the main activity of Echoes of Havenbay change from investigating primarily fresh murders to reopening cases that Havenbay considers closed, solved, cold, misclassified, or no longer worth asking about?

Provisional answer

Yes, subject to a falsification test and a narrower rule than either “cold cases” or “solved cases.”

The strongest synthesis is:

Internal story engine: The record is wrong, and somebody is still living with the consequence.

Player-facing promise: Ally Quinn reopens the cases Havenbay has decided are over.

Episode construction rule: A contested record, a live human consequence, reconstructive proof, and an accountable resolution.

This is a better backbone than a default series of fresh manhunts. It aligns Ally’s lawful role, the merge mechanic, the title Echoes of Havenbay, and the Thomas Quinn season arc.

It is not yet safe to make the backbone exclusive. A historical error is not automatically present-tense jeopardy. A closed case is not automatically public in full. A passive killer is not automatically an antagonist. Those claims need correction before canon changes.

Authority and limits

This brief uses the current versions of:

1. pivot-brief-2026-08-23.md, current-state authority.
2. critical-reviewer-charter.md, assessment authority.
3. episode-1-vo-lock-v1.1-WIP.md, prose and listening-speed authority.
4. ally-quinn-bible.md, except where superseded.
5. canon-character-roster.md and portrait-prompts-canon-cast.md, character authority.

The three stale art-pipeline documents are excluded.

This brief does not have the full text of Slack Water or the current season straw man. It cannot judge that pitch or specify exact season placement without those documents. It also cannot establish what the twelve original testers meant by “no jeopardy” from the phrase alone.

Evidence that governs the decision

- The Listener failed twelve of twelve friends-and-family tests. Reported reasons included no jeopardy, too safe, not real true crime, and non-completion.
- The current pivot requires a murder-led pilot, credible present-tense jeopardy, four or more material decisions, and a pause-safe clock.
- Internal numeric scoring is abolished as decision input. Forced ranking and direct tester evidence take precedence.
- The game activity is reconstruction. The player obtains and physically combines recordings, records, manifests, images, and related residue into auditable conclusions.
- Ally is a podcaster and private investigator. Her licence does not grant police powers or automatic access to protected evidence.
- Thomas Quinn’s death is itself an officially settled case whose record may be wrong.
- The bible currently uses “cold cases” for short between-episode repeatables. If the main episodes adopt this backbone, that terminology must change.

Where the two positions already agree

Claude’s formulation identifies the central thematic and mechanical opportunity correctly.

Point	Agreement	Why it matters
Series engine	Strong agreement	“The record is wrong, and somebody is still living with the consequence” contains both investigation and human cost.
Title fit	Strong agreement	An echo is a past act still sounding in the present. The title supports retrospective cases better than routine pursuit.
Player activity	Strong agreement	Reconstructing a record maps naturally to merge play. A manhunt is harder to express honestly on a board.
Ally’s role	Strong agreement	Public proceedings, family-held documents, archives, interviews, and voluntary disclosures give Ally a credible starting position.
Season integration	Strong agreement	Thomas becomes the largest version of Ally’s weekly practice, not a separate mythology task.
Variety need	Strong agreement	The record-error shape must vary. Wrongful conviction cannot become the default formula.
The Listener	Strong agreement	Its demotion is correct. It still needs a dread and jeopardy pass before its live-case status can carry it.
Testing before canon	Strong agreement	The remaining dispute depends on what players experience as jeopardy, not on either model’s confidence.

Claims that need amendment

1. “A closed case is public”

This is directionally useful and literally too broad.

Some judgments, hearing transcripts, inquest findings, press archives, public registers, and family-owned records may be available. Police files, exhibits, sealed material, medical records, protected witness information, and retained or destroyed evidence may not be. Jurisdiction and case history matter.

Consensus correction: Access improves materially, but Kill Gate Zero remains. Every pitch must name:

- who lawfully possesses each decisive item;
- why that person can disclose it;
- what remains inaccessible;
- how Ally proceeds without police privilege.

2. “The antagonist does not need to act”

This solves one kind of contrivance but risks removing dramatic opposition.

The wrongness of a record creates mystery and injustice. It does not automatically create jeopardy. A free killer who quietly benefits from an old verdict may be a subject, but not necessarily an active antagonist. A wrongfully imprisoned person is suffering current harm, yet Ally may still lack an opposing force unless somebody has a credible reason and lawful means to resist correction.

Consensus correction: The antagonist need not hunt Ally. The episode does need an active resistance system. It may be a person, institution, family interest, reputational economy, evidentiary constraint, or deadline. Its behaviour must be rational under scrutiny.

3. “The clock problem ends”

Closed cases support better pause-safe clocks, but they do not remove the problem.

A hearing, demolition, evidence-disposal date, scheduled transfer, or imminent public decision can work. Repeated dying witnesses, sudden demolitions, and arbitrary deadlines will become visible machinery. A limitation period is not a universal murder clock and must not be used without jurisdiction-specific support.

Consensus correction: A clock is optional when the current consequence is already irreversible and accumulating. If used, it must arise from the case rather than from the need to manufacture urgency. The clock convention in the pivot brief remains authoritative.

4. “A killer walking around Havenbay creates permanent jeopardy”

Not by itself. Freedom is an unresolved moral and public-safety consequence. Present jeopardy requires a credible current exposure: continued access to a victim type, coercion of a keeper, evidence destruction, retaliation, institutional power, or a live decision that will harden the false record.

Consensus correction: Historical murder plus present impunity is a foundation. The pitch must identify who can still be harmed now, by what mechanism, and why the harm is credible.

Proposed shared backbone

The backbone should be adopted as a default franchise promise, with controlled exceptions for fresh cases.

Canon wording candidate

Echoes of Havenbay follows podcaster and private investigator Ally Quinn as she re-opens cases the city calls closed, solved, cold, or no longer worth asking about. In every main episode, the official record is contested and somebody is still living with the consequence. The player reconstructs what happened from lawfully obtained residue, makes material investigative choices, and decides what account Ally can responsibly put into the world.

Four-part episode engine

Every main episode must contain all four elements:

1. **Contested record** - an official, accepted, or socially dominant account that can be stated plainly.
2. **Live consequence** - a specific person is still paying, benefiting, being exposed, or being denied remedy now.
3. **Reconstructive proof** - the conclusion is built from obtained evidence and can be audited by the player.
4. **Accountable resolution** - the player changes what Ally does, what evidence becomes available, or what ending is reached. Resolution does not require a conviction, but it must answer the dramatic question and reckon with publication.

Historical injustice alone fails. Atmosphere alone fails. A twist without traceable proof fails. A final dialogue-only branch fails.

Fresh cases remain possible

Fresh murders should be exceptions, not forbidden material. They can provide season variation and maximum immediacy. They must still fit the same reconstructive and ethical engine. This preserves the possibility that the best pilot is fresh while making wrongly settled cases the series identity.

Two-axis case taxonomy

Claude's five record shapes are useful. They need a second axis so variety is structural rather than cosmetic.

Axis one: What is wrong with the record?

Record shape	Core question
Never solved	Why did the inquiry stop before the answer?
Wrongful conviction	Who benefited from certainty about the wrong person?
Misclassified	What was called accident, suicide, disappearance, or natural death instead of murder?
Covered up	Which institution controlled the surviving account?
Officially solved but contested	Which apparently decisive fact does not withstand reconstruction?

Wrongful conviction should provisionally be limited to one main episode per season until a blind shelf test establishes that repetition does not feel formulaic.

Axis two: What is happening now?

Live consequence	Present dramatic pressure
Wrongful loss of liberty or status	A person remains confined, excluded, or publicly condemned.
Continued access or power	The responsible person can still harm, coerce, profit, or select victims.
Keeper under pressure	A survivor or witness is still being intimidated, disbelieved, displaced, or exploited.
Evidence opportunity closing	A lawful, non-arbitrary chance to preserve or test evidence will end.
Institutional decision	A hearing, transfer, demolition, appointment, payment, or official declaration will harden the outcome.
Public lie blocking remedy	The accepted story prevents recognition, compensation, protection, or mourning.

A pitch should pair one record shape with one or more live consequences. Repeating either axis too closely risks sameness.

Proposed jeopardy rule for adjudication

The current two-track rule may itself be forcing contrived threats to Ally. It should not be silently discarded, but it should be tested against a replacement.

Candidate replacement

Jeopardy gate: The episode must place a named person under a credible current consequence and must put something Ally is responsible for at risk through investigation or publication. Physical danger is one valid form, not the default requirement. Historical suffering without a live mechanism fails. Professional embarrassment without meaningful cost fails. The weaker track controls the read.

The two tracks become:

1. **Human consequence:** What can still happen to the keeper, accused person, survivor, future victim, or community because the record remains wrong?
2. **Investigative consequence:** What can Ally's choices preserve, expose, worsen, lose, or make irreversible?

For the pilot, the threshold should remain higher than for later episodes. It must be murder-led and the live consequence must be legible in the cold open. A buried fact becoming interesting is not enough.

Cheapest useful test

Claude's proposed question is fast but leading. It asks testers to predict their preference using the desired categories. Stated preference is weaker than an observed choice.

Recommended falsification test

Run a blind six-card shelf test with at least five previous testers, ideally all twelve.

- Write six equal-length, title-free episode cards.
- Include two fresh murders, two genuine cold cases, and two officially solved but contested cases.
- Give every card a named keeper, a concrete present consequence, one reconstructive activity, and comparable specificity.
- Remove category labels and franchise explanations.
- Ask each tester to choose the story they would start first, the story they would try second, and the story they would skip.
- After each choice, ask only: "What made that feel urgent or safe?"
- Then show the backbone sentence and ask whether it makes the series more or less attractive.

Decision rule without scores

- **Adopt the backbone** if contested-record stories are repeatedly chosen first or second for reasons tied to injustice, consequence, discovery, or reconstruction.
- **Reject exclusivity** if fresh murders dominate first choices because current physical danger is the essential hook.
- **Reject the current copy, not necessarily the backbone** if testers like the category only after it is named. That indicates the cards failed to dramatise the live consequence.
- **Do not treat a split as a tie.** Force-rank the six cards and state the dominant reason. If the sample remains divided, preserve a mixed season and test a playable cold open from each family.

After ranking, Claude's binary question can be used as a diagnostic follow-up, not as the decision test.

Implications for the current slate

- The Listener remains demoted. Making fresh cases rarer does not repair its existing failure.
- Slack Water may fit the proposed identity, but this brief cannot rank or approve it without the full pitch.

- Any current pitch must be re-read for lawful evidence custody, rational resistance, present consequence, and the player's four or more material decisions. A new backbone does not grandfather weak causality.
- The best pilot is the story that proves the franchise promise at listening speed. A pure statement of theme that opens slowly can still be the wrong pilot.

Document impact if adopted

The change reaches further than five isolated edits. At minimum:

1. Pivot brief: high concept, pilot doctrine, present-tense jeopardy, clock examples, and case-family language.
2. Critical reviewer charter: matching gates and fact-versus-inference checks.
3. Story-generation prompt: version 2.1 or later, including evidence custody and two-axis taxonomy.
4. Ally Quinn bible: high concept, Ally's practice, pilot doctrine, season architecture, and the superseded twelve-lead material.
5. Bible section 6.4: rename or redefine the current between-episode "cold cases" to prevent a terminology collision.
6. Season straw man: forced ranking of episode families and deliberate variation.
7. Pitch framework: current consequence, access map, resistance logic, and player decisions become mandatory fields.
8. Marketing promise: confirm whether "reopening the cases Havenbay calls over" is the clearest player-facing proposition.

No canon file should be changed until the blind test is run or the creator explicitly accepts the remaining risk.

Questions Claude must adjudicate

Claude should answer each item with **AGREE**, **DISAGREE**, or **AMEND**, then supply exact replacement wording.

1. Is the proposed three-level wording the right separation between internal engine, player promise, and episode rule?
2. Does access improve rather than become automatic, and should evidence custody remain a kill gate?
3. Can a passive beneficiary of a wrong record provide enough opposition, or must every episode contain an active resistance system?
4. Does the current two-track jeopardy rule need the proposed replacement?
5. Should fresh cases remain controlled exceptions?
6. Is the two-axis taxonomy sufficient to prevent sameness?
7. Is the blind six-card test a stronger falsifier than the proposed binary preference question?
8. Which assertion in this brief is weakest, and what evidence would overturn it?
9. What cannot be judged without the full Slack Water pitch and current season straw man?
10. What exact consensus wording should be presented to the creator after this round?

Consensus protocol

Absolute agreement should mean explicit convergence, not mutual politeness.

For each round:

1. Preserve a table of settled points, amended points, and unresolved points.
2. Label every material claim as document fact, tester evidence, general inference, or creative preference.
3. State the strongest counterargument to the preferred answer.
4. Do not use scores or percentages.
5. Do not infer missing pitch details.
6. Do not edit canon while a governing rule is still disputed.

7. End with one of three verdicts: **consensus reached**, **provisional consensus pending test**, or **material disagreement remains**.

The present verdict is **provisional consensus pending test**.

Revision history

Version	Date	Change
1.0	24 August 2026	First cross-model consensus brief. Separates agreement from unresolved claims, proposes a two-axis taxonomy, revised jeopardy gate, and blind shelf test.