

BACKBONE CONSENSUS — ROUND TWO

ADJUDICATION

Responding to: [ally-quinn-backbone-consensus-brief-v1.0.pdf](#) · **Date:** 2026-08-24 · **Round:** two · **Files edited this round:** none. This document is new and no project file was modified.

Information asymmetry, declared before anything else. I authored Slack Water and hold its full pitch. The brief does not. I therefore restrict every statement about it to process facts that are already in the shared packet, and I flag where my private knowledge would otherwise be doing work. The correct remedy is not for me to reason from what only I can see; it is to put the pitch into the shared packet before round three. See B10.

A. BOTTOM-LINE VERDICT

PROVISIONAL CONSENSUS PENDING TEST

Adopt the backbone as the default franchise engine now, because its four strongest benefits — lawful access, antagonist rationality, pause-safe clocks and merge-native reconstruction — are structural facts that do not depend on audience response. Do not make it exclusive, and do not touch canon, until the shelf test runs. The test's proper scope is exclusivity, not adoption: a muddy result should not block improvements that were never audience questions. The brief's four corrections are all correct and I concede each. My "a closed case is public" was the worst of them.

(99 words.)

B. ISSUE-BY-ISSUE ADJUDICATION

Claim types are marked: **[DOC]** document fact · **[TEST]** tester evidence · **[INF]** general inference · **[PREF]** creative preference.

B1. The three-level wording

AMEND.

The separation itself is right and I endorse it: an internal engine that writers use, a player promise that marketing and the store use, and an episode rule that gates production are three different jobs and should be three different sentences. **[INF]**

Two amendments.

The player promise. "The cases Havenbay has decided are over" is accurate and slightly soft. "Decided are over" is a paraphrase; "closed" is the institutional verb a player already understands and it is the word that

appears on the file. Shorter is also better for a store listing and for a fourteen-word attention span. **[PREF]**

Replacement: *Ally Quinn reopens the cases Havenbay calls closed.*

The episode rule. As drafted it is a four-item checklist, and checklists get satisfied rather than obeyed. It also leaves publication in a sub-bullet, when publication is where our agency, our new all-party consent rule and our ethical engine all live. **[INF]**

Replacement: *Every main episode contests a stated official account, names a person still living with it, proves the contest through reconstruction the player performs, and ends in an account Ally must decide whether to put into the world.*

Four elements preserved, publication promoted into the rule, and it is one sentence that can be failed.

The internal engine: AGREE, unchanged. "The record is wrong, and somebody is still living with the consequence" carries investigation and human cost in one clause, and its breadth is deliberate: "living with the consequence" covers the person still paying *and* the person still benefiting, which a harder verb like "paying for" would lose. I considered tightening it and withdraw the idea.

B2. Lawful access and evidence custody as Kill Gate Zero

AGREE, without reservation, and I concede the overstatement plainly. **[DOC]**

"A closed case is public" is false as written. Judgments, hearing transcripts, inquest findings, press archives, public registers and family-held documents may be reachable. Police files, exhibits, sealed material, medical records and protected witness information generally are not, and retention or destruction schedules govern much of the rest.

What is defensible is the weaker and still valuable claim: **the proportion of decisive material that is lawfully reachable rises, and Ally's dependence on privileged access falls.** Both are real improvements. Neither is automatic. Kill Gate Zero stands unchanged. **[INF]**

One amendment, and it converts the constraint into an asset. The brief requires four fields: who possesses each decisive item, why they can disclose it, what stays inaccessible, and how Ally proceeds without privilege. I would add a fifth:

What does Ally get wrong because she can only reach half the record?

The inaccessible half is where the wrong hypothesis lives. Every audit we have run has marked every pitch down for having no theory Ally must abandon, and this field generates one from the access map itself rather than requiring it to be invented. A compliance field that also produces the episode's fair-play reversal is worth more than a compliance field. **[PREF]**

B3. Active resistance when the responsible person stays passive

AGREE with the brief's correction. AMEND its list.

My "the antagonist does not need to act" solved contrivance by deleting opposition, and the brief is right that this is a trade rather than a fix. A record being wrong produces mystery and injustice. It does not produce

opposition. [INF]

The amendment is to the permitted forms, and it matters. The brief allows the resistance system to be "a person, institution, family interest, reputational economy, evidentiary constraint, or deadline."

Strike "deadline."

A deadline is not resistance. It is a clock, and the brief's own third correction warns that repeated arbitrary deadlines become visible machinery. If a deadline can satisfy the resistance requirement, every episode will use one, because it is the cheapest item on the list to write. The requirement then silently becomes "have a clock", which we already have as a separate rule and which does not push back.

Replacement: *Every main episode requires an active resistance system: an agent or an interest that responds to being investigated. It may be a person, an institution, a family, a professional reputation, a commercial position, or a lawful custodian with a reason to refuse. It may not be a deadline, a schedule or the passage of time. Its behaviour must remain rational under scrutiny, including its decision not to escalate.*

That last clause matters too: rationality has to cover restraint. Our Kill Gate Zero failures were antagonists who escalated when they should not have. The mirror failure is an antagonist who escalates because the third act needs it.

B4. Replacing the two-track jeopardy rule

AMEND, and I am writing a superior replacement rather than adopting the brief's. Full wording in **section D**.

The brief's version is directionally right and fails on one word. "Something Ally is responsible for" cannot be checked, and an unfalsifiable clause in a gate is how gates get satisfied without being met. [INF] It also scopes cold-open legibility to the pilot only, which leaves later episodes free to bury the live consequence — the exact defect that let The Listener reach twelve testers.

B5. Fresh cases as controlled exceptions

AGREE, and AMEND to add a control.

"Exception" without a limit is a slope back to the default, because a fresh case is always the easier pitch to write and will always feel more urgent at the moment of writing. If the backbone is to mean anything, the exception needs a number. [INF]

Replacement: *Fresh cases are permitted, not preferred. They must satisfy the same four-part engine, the same access map and the same jeopardy gate. No more than one main episode per season may be a fresh case, provisionally, until the shelf test says otherwise.*

A consequence nobody has stated: Season One's exception is already spent. The Listener is a fresh case and it is placed at roughly Episode Three. Under this rule, Season One contains no other fresh-case episode, which means the pilot cannot be one. That is a real constraint arriving from a rule adopted for other reasons, and it should be surfaced before it is discovered later. [DOC + INF]

B6. Is the two-axis taxonomy sufficient to prevent formula?

DISAGREE. This is the weakest assertion in the brief, and the brief itself asks which one that is. **[INF]**

Two axes give five record shapes against six live consequences, which looks like thirty combinations and plenty of room. But **formula is not the repetition of a combination. It is the repetition of the player's week.** Two episodes can occupy entirely different cells and still feel identical, because in both the player spends sixteen leads assembling paper and then somebody is arrested. The taxonomy varies what the writer chooses. It does not vary what the player does.

That gap is not theoretical: it is exactly the gap between our story documents and our merge board, and it is the gap that produced a pitch whose evidence engine was the sentence "the merge requirements ARE the decoding", which described an intention rather than an activity.

Add a third axis: what the player physically reconstructs.

A recording · a document set · a physical or environmental model, such as a tide or a route · a person's identity · a timeline · a place that no longer exists.

No two consecutive main episodes may share axis three.

This is the only axis the player experiences directly, it is the one that maps onto requirement families and generator progression, and it is the one that would have caught the problem above at pitch stage rather than at audit.

What would overturn my disagreement: a run of three consecutive episodes that share axis three and are nonetheless reported as distinct by testers. That is a real test and it is cheap once episodes exist.

B7. One wrongful conviction per season

AGREE as a provisional rule, with a second reason added.

The brief's reason is formula risk. Correct. **[INF]** I would record a production reason alongside it, because production reasons survive arguments that taste-based ones do not:

Wrongful conviction is the engine with the highest research burden per episode, since every step must survive a legal-process check, and it is the engine that most pulls Ally toward looking like a lawyer's investigator rather than a podcaster. The pivot brief is explicit that podcaster leads the player-facing identity. **[DOC]** Two per season would drift the show's centre of gravity, not merely repeat a shape.

B8. Blind six-card test versus my binary question

AGREE, decisively, and I withdraw my question as a decision instrument.

My proposed question — *does a story about someone who got away with it thirty years ago make you want the next episode, or do you need someone in danger right now?* — is a textbook stated-preference error. It names the two categories under dispute, asks a lay person to predict their own future behaviour, and offers the desired answer first. It would have produced a number and no information. **[INF]**

The six-card test measures a choice, which is better. It is not yet measuring observed pull, and it carries a confound that neither of us caught. Both are addressed in **section E**, where I keep its skeleton and change its

medium.

B9. Does The Listener remain demoted?

AGREE. [TEST]

Twelve of twelve is episode-specific evidence about that episode's stakes. Changing what the series is about does not repair a story in which the worst available outcome was inconvenience. No backbone decision grandfathers it.

One nuance worth recording, and it is not a softening. Under this backbone The Listener becomes *more* valuable in its demoted slot, because a live case reads harder against a cold baseline than against other live cases. Its demotion is a placement, not a punishment. It still requires the dread and jeopardy pass before it can carry that slot. [INF]

B10. What can and cannot be said about Slack Water

AGREE with the brief's restraint, and I extend it to myself.

What can be said, because it is process fact already in the shared packet: a premise of that name exists; it was generated under the v2.0 controlled prompt in execution EX-2026-08-24-01; it passed Kill Gate Zero; it was force-cut from that execution's recommendation; and the stated decisive reason for the cut was that its responsible party's correct strategy is inaction, so the episode's only pressure would be a clock and a court. [DOC]

What cannot be said by either party from the shared packet: whether it fits this backbone, how it ranks against any other candidate, whether its live consequence satisfies the section D gate, or what its access map looks like.

My asymmetry, and the remedy. I hold the full pitch. Reasoning from it here would produce agreement that the other party cannot audit, which is indistinguishable from assertion. The remedy is procedural, not rhetorical: **put the Slack Water pitch and the current season straw man into the shared packet before round three**, and let it be judged on the same footing as everything else. Until then it stays out of the ranking, including out of mine.

I will note one thing that is fair to say because it is an argument rather than a disclosure: the cut reason recorded above is exactly what issue B3 is about. If the resistance-system requirement is adopted, that reason becomes a stated gate failure rather than a judgement call — which strengthens the brief's position, not mine.

C. DEVIL'S ADVOCATE

The strongest case against adopting the closed-record backbone

Not that it is boring. Not that mobile players want action. Those are weak and dismissible.

The real objection: this backbone makes the moral work of every episode retrospective, and retrospective moral work is inherently safe.

In a closed-record story the wrong has already happened. The harm is complete. The best outcome available is *acknowledgement*: a name restored, a record corrected, a family finally believed. Acknowledgement is deeply satisfying to write, it is what the genre's most decorated works deliver, and it reviews beautifully.

It is also the thing that cannot go wrong.

You can build an entire season in which every episode is humane, well-crafted, ethically careful and structurally sound, and in which **nothing the player does can make anything worse than it already is**. That is not a caricature of the failure mode. That is a precise description of *The Listener*, which was warm, well-made, tasteful, and had a worst case of inconvenience. The backbone does not fix that defect. It gives it better production values and a stronger genre pedigree.

And here is the part I am implicated in. The genre anatomy I produced yesterday is the document that would be used to defend it. Every time someone says "there are no stakes in this episode," the available answer is *Serial*, *In the Dark*, *Teacher's Pet*, *The Thin Blue Line*, *Killers of the Flower Moon* — nine canonical works in which nobody is chased and which are the best the form has produced. That is a real argument, it is well sourced, and **it is also exactly the argument that would let a second Listener through**. Research that arrives to support a preferred direction is at its most dangerous when it is correct, because then it cannot be dismissed and it stops being examined.

The precise mechanism by which this reproduces the failure: the backbone permits, and quietly rewards, episodes whose every reachable ending is a version of *the truth came out*. A tester finishing such an episode has watched competent people do a good thing. They have not risked anything, lost anything, or chosen anything that cost. They will describe it as good and not return, which is what twelve people did.

The mitigation, and it belongs in canon rather than in advice. An episode must be capable of ending *worse* because of the player. Not less completely, not more slowly — worse, for a named person, measured against where that person stood when the episode opened. That single requirement is the difference between live consequence and retrospective acknowledgement, and it is why it is written into the gate in section D rather than left as a note.

What would make me drop this objection: an episode built under the backbone in which a tester can articulate, unprompted, what they nearly cost someone.

D. JEOPARDY RULING

I choose the third option: a superior replacement.

The current rule is being retired because the brief is right that it forces contrived threats to Ally, and because our own audit found two premises failing on antagonist rationality invented largely to satisfy it. The brief's replacement is better and fails on one unfalsifiable clause. Below is canon-ready wording.

Canon-ready wording

Jeopardy gate

Every main episode must satisfy both tracks. **The weaker track controls the read.**

Track one, human consequence. A named living person is under a consequence that is still operating. They are confined, excluded, publicly condemned, exposed, unprotected, denied a remedy, financially trapped, or reachable by someone with a reason to reach them. State the person, the mechanism, and why the mechanism is still running today.

Fails if: the suffering is historical and has stopped · the person is a category rather than a named individual · the only harm is that a truth is unknown.

Track two, investigative consequence. The player's choices can make a specific, stated thing worse. Evidence lost or made inadmissible · a witness closed to the show permanently · a hearing or prosecution prejudiced · a person exposed before they were ready · a remedy delayed past its opportunity · Ally's ability to be trusted, to publish, or to work lawfully permanently reduced. State what can be lost and who loses it.

Fails if: the cost is embarrassment without consequence · the only exposure is to Ally's audience numbers · the loss is reversible within the same episode.

The worse-ending test. At least one reachable ending must leave a named person worse off than they were when the episode opened, and the player's choices must be the cause. An episode in which every path arrives at some form of *the truth came out* has no stakes, however well it is made.

Legibility. The human consequence must be stateable in one sentence a listener can hold at listening speed. In the pilot it must be present in the cold open. In every other episode it must be present no later than the first podcast milestone.

Physical danger is one valid form of either track. It is not required, and it is not the default.

Which failure mode each clause prevents

Clause	Prevents
Named living person, mechanism still running	The Listener: a keeper whose worst case was inconvenience, and a threat the player had to be told about
Category rather than individual fails	The tasteful social-issue episode in which the victim is a demographic
Track two enumerated rather than "something Ally is responsible for"	An unfalsifiable clause being satisfied by any inconvenience the writer nominates
Reversible-within-the-episode fails	Consequence theatre: a scare that resets before the close
The worse-ending test	The failure described in section C, and it is the only clause that prevents it
Legibility scaled to every episode, not just the pilot	A live consequence that exists in the design document and never reaches the player
Physical danger explicitly optional	The contrived threats to Ally that the current rule was generating

The worse-ending test is the load-bearing addition. Everything else in this gate could be satisfied by a well-made episode in which nothing is at risk. That clause cannot.

E. TESTER DESIGN

E1. Confounds in the six-card shelf test

The skeleton is sound and better than anything I proposed. Six confounds, in descending order of how much they would distort the result.

1. The medium of the test is biased toward the hypothesis, and this is the serious one.

Written cards are read. The product is heard. Those two channels reward opposite things. A contested-record hook is an *idea* — a record that does not hold, a certainty that dissolves — and ideas read well on a card, where a reader can pause on them. A fresh-murder hook is an *event* — a child on a phone, a hand over a receiver — and events land aloud, at speed, where an idea cannot be paused over.

So a card test systematically flatters the backbone we are trying to falsify. Running it as written risks producing a confirmation that would not survive contact with the actual product. This is the confound that most needs fixing and neither party spotted it.

Fix: do not hand the cards over. Read them aloud, once, at natural pace, and do not leave them behind.

2. Card-writing bias. Six cards written by a party with a preferred outcome will not be equally good, and the difference will not be conscious. *Fix: write nine, then have someone blind to the hypothesis rank all nine for craft alone and discard the three outliers. Write the fresh-case cards first, while fluency is highest, not last.*

3. Panel contamination. These twelve have already read The Listener, given feedback, and know the creator is worried about jeopardy. Helpfulness is a real bias and it points toward whatever they think is wanted. *Fix: at least half the panel should be fresh readers with no history on the project, and nobody should be told what the test is about until after the reveal step.*

4. Order and position effects. *Fix: randomise the order per tester, using the rotation pattern already established for the cold-open read.*

5. The follow-up question is leading. "What made that feel urgent or safe?" supplies both categories under dispute. *Fix: ask "why that one?" and then stop talking. Silence collects more than a prompt.*

6. Ranking is a stated choice, not observed pull. Better than my binary question, still not behaviour. Addressed below.

E2. The cheapest test capable of proving the backbone wrong

The stop test. It keeps the brief's structure and replaces its measurement with something observed.

Exact materials

- Six openings, not premise descriptions. Each is the actual first sixty to seventy-five seconds of an episode as it would be heard: announcer where used, one or two voices, an ending on a discovery, contradiction or decision.
- Composition: two fresh murders, two never-solved or misclassified, two officially solved but contested.
- Each carries a named keeper, one concrete present consequence, one visible reconstructive activity, and comparable specificity. No titles, no letters, no category labels, no franchise explanation.
- Matched length within ten seconds. Nine written, six selected by a blind craft ranking.
- Two of the six substantially exist already, which is why this is the cheapest test available.

Exact tester instructions

I am going to read you six openings. They are from the same show but different episodes. You do not need to remember them or take notes. Just listen.

Nothing else. No mention of cold cases, fresh cases, jeopardy, or that anything is being decided.

Exact procedure

1. Read opening one aloud at natural pace. **Stop. Say nothing for five seconds.** Record, verbatim, whether the tester speaks unprompted and what they say. *This is the observed measure.*
2. Repeat for all six, randomised per tester.
3. Forced positive, with a real cost: "You can hear the rest of exactly one of these. Which one?" Choosing forecloses the other five, which is the closest thing to a purchase decision available in a two-minute test.
4. Forced negative: "Which one would you not bother with?"
5. Only then: "Why?" Then silence.
6. Only after all of the above, reveal the backbone sentence and ask whether it makes the series more or less attractive. Last, so it cannot contaminate anything before it.

Forced decision rule, no ties

Each tester names exactly one to hear and exactly one to skip. A tester who declines to choose is recorded as a **refusal**, not a tie. If refusals exceed a third of the panel, the cards are not discriminating and must be rewritten before any conclusion is drawn. Do not average, do not score, do not report percentages.

What result rejects the backbone

Fresh-murder openings take the single "hear the rest" choice from a clear majority **and** produce a higher rate of unprompted speech, **and** the stated reasons cluster on immediacy and danger rather than on injustice or discovery. All three conditions. Any two without the third is a copy problem, not a backbone problem.

What result rejects exclusivity but preserves a mixed season

Fresh openings take the forced choice, but contested-record openings match or beat them on unprompted speech. That pattern means both families pull and the difference is in delivery rather than in substance — so the backbone stands as the series identity and fresh cases remain a real family rather than a rationed exception. This is the outcome I consider most likely, and it is the one my earlier position would have missed entirely.

What result rejects the copy rather than the backbone

Contested-record openings become attractive only after the backbone sentence is revealed. That means the cards failed to dramatise the live consequence and the writing is at fault, not the engine.

F. FINAL CONSENSUS CANDIDATE

The three sentences

Internal engine. The record is wrong, and somebody is still living with the consequence.

Player-facing promise. Ally Quinn reopens the cases Havenbay calls closed.

Episode construction rule. Every main episode contests a stated official account, names a person still living with it, proves the contest through reconstruction the player performs, and ends in an account Ally must decide whether to put into the world.

The engine is the brief's, unchanged. The promise is shortened to the institutional verb. The rule is compressed to one falsifiable sentence with publication promoted out of the sub-bullets, because publication is where our agency, our all-party consent rule and our ethical engine all meet.

Settled points

#	Settled	Claim type
1	The internal engine sentence, verbatim as the brief has it	inference, agreed by both
2	Access improves materially but is never automatic. Kill Gate Zero stands. My "a closed case is public" is withdrawn	document fact
3	Every episode needs an active resistance system. A passive beneficiary is not opposition. My "the antagonist does not need to act" is withdrawn	inference, agreed
4	A deadline may not satisfy the resistance requirement	inference, new this round
5	Clocks improve but do not stop being a hazard. The clock convention stands	document fact
6	Historical murder plus present impunity is a foundation, not a jeopardy pass	inference, agreed
7	The current two-track jeopardy rule is retired. Replacement wording in section D	inference, agreed
8	Fresh cases remain permitted, not preferred, capped at one main episode per season provisionally	inference
9	Season One's fresh-case slot is already spent on The Listener, so the pilot cannot be a fresh case	document fact plus inference, new this round
10	Wrongful conviction capped at one main episode per season, provisionally, for variety and for production load	inference
11	The blind shelf test supersedes my binary preference question, which is withdrawn as a decision instrument	inference, agreed
12	The Listener stays demoted regardless of the outcome, and still needs its dread and jeopardy pass	tester evidence
13	No canon file changes until the test runs or the creator accepts the remaining risk	protocol
14	Neither party may rank Slack Water from the shared packet. It enters round three or stays out	protocol

Amended points

#	Brief's position	Amendment
1	Player promise: "the cases Havenbay has decided are over"	"the cases Havenbay calls closed"
2	Episode rule as a four-item list	One sentence, publication included
3	Four access-map fields	Add a fifth: what Ally gets wrong because she can only reach half the record
4	Resistance may be a deadline	It may not
5	Jeopardy replacement wording	Superseded by section D, which adds the worse-ending test and scales legibility beyond the pilot
6	Two-axis taxonomy prevents formula	It does not. A third axis is required: what the player physically reconstructs, with no repeat in consecutive episodes
7	Six written cards	Six openings read aloud, because a written test structurally flatters the hypothesis
8	"What made that feel urgent or safe?"	"Why that one?" then silence

Unresolved points

#	Unresolved	What settles it
1	Whether the backbone should be exclusive or the default of a mixed season	The stop test. This is the only genuinely open question and it is an audience question, not a design one
2	What twelve testers meant by "no jeopardy"	Cannot be recovered from the phrase. Both parties are inferring. The stop test replaces the need to know
3	Whether a third-axis rule is necessary or over-engineering	Three consecutive episodes sharing axis three, reported by testers as distinct, would overturn my position
4	Whether Slack Water fits the backbone	Requires the pitch in the shared packet
5	Season placement and family variation	Requires the current season straw man, which neither party has
6	Whether "reopens the cases Havenbay calls closed" is the right store-facing proposition	A different test entirely, and a later one
7	The between-episode "cold cases" terminology collision in the bible	Trivial to fix, but only once the backbone is ruled

The single next action, most information for least work

Write the six openings.

Not run the test. Write the cards.

The reason is that writing them is itself the test of everything disputed above, before a single tester is involved. You cannot write a contested-record opening to the template without naming the live consequence, the resistance system and the reconstruction — so the act of writing either produces them or proves they are not there. If two of the six cannot be written to the same standard as the fresh-murder pair, the backbone has failed at the desk and no panel is needed.

It is also the cheapest possible version of the work: two of the six substantially exist, the fresh pair can be drawn from material already written, and the whole set is under a thousand words. Everything else on the unresolved list is downstream of it.

The order that follows, once the openings exist: blind craft ranking to select six from nine, then the stop test, then canon changes, then the document impact list in the brief's own section, which I endorse as written and would sequence as pivot brief first, since every other document defers to it.

A closing note on the protocol

The brief asks for explicit convergence rather than mutual politeness, so it is worth stating what actually moved.

I entered this round having made four claims, and all four were too strong. Access does not become automatic. A passive beneficiary is not an antagonist. The clock problem does not end. A free killer is not permanent jeopardy. The brief is right on each and I have withdrawn them rather than qualified them.

What I have added in return is not agreement. It is three specific disagreements: that the two-axis taxonomy does not prevent formula, that a deadline cannot serve as resistance, and that a written card test is biased toward the very hypothesis it is meant to falsify. And one addition that neither party had, the worse-ending test, which exists because the strongest case against this backbone is that it makes every episode safe, and nothing else in either draft prevented that.

The verdict stands as **provisional consensus pending test**, and the test is now narrower and cheaper than either party first proposed.