

ALLY QUINN STORY BACKBONE —

CONSENSUS v1.2

Document ID: AQ-NAR-BACKBONE-CONSENSUS · **Version:** 1.2 · **Date:** 2026-08-24 · **Status:** PROVISIONAL CONSENSUS, ONE WORDING ISSUE OPEN · **Supersedes:** v1.0, v1.1, and `backbone-consensus-round2-adjudication.md`

Canon impact: none yet. No project file has been edited in any round of this process.

PART ONE — ROUND THREE ADJUDICATION

A. Verdict

PROVISIONAL CONSENSUS — ONE WORDING ISSUE REMAINS

The remaining issue is narrow and I supply exact final wording for it in section C. If that wording is accepted, this process reaches consensus and no further round is needed on the rule.

Separately, and not part of the rule dispute: the six-opening experiment is **not yet ready to execute**, for four specific reasons given in section E, one of which is a regression from the previous round rather than an outstanding item.

B. Rule-by-rule adjudication

B1. Structural tools adopted now; the spoken-opening test decides dominance and pilot family rather than permission.

AGREE. The three-decision scope is better than my two-way split and I adopt it. A backbone is not only a production technique; it is a repeated audience promise and a season allocation, and tester behaviour is properly relevant to both.

I also accept the precision correction attached to it. I listed antagonist rationality among the benefits that do not depend on audience response, which over-claimed. What an old case does is **remove the need for a killer to stage elaborate present-day crimes**, which eliminates one failure mode. It does not supply opposition. Opposition still has to be written, every time. That is a reduction in failure surface, not a structural gift, and my earlier phrasing blurred the two.

B2. The test does not decide exclusivity, because exclusivity is not recommended.

AGREE. Exclusivity was my word, introduced in round two and never proposed by anyone. Withdrawn.

B3. A third axis, what the player reconstructs, is mandatory at pitch and season level.

AGREE, and the expanded candidate list in v1.1 is better than mine. Custody history of an object and network of payments are both stronger merge subjects than two of my originals.

B4. Resolution mode must also be tracked.

AGREE, and this is a correction to me rather than an addition to me. I claimed axis three was "the only axis the player experiences directly", which is plainly too broad: endings are experienced directly, and so are decisions and pacing. Worse, the example that defeats my claim is my own argument turned around. "Sixteen leads of paper, then an arrest" is caught by the **pair** of reconstruction mode and resolution mode. Reconstruction mode alone catches only the first half of the sentence.

B5. Resistance and clock remain separate pitch fields.

AGREE, unchanged.

B6. No one-fresh-case quota was proposed.

AGREE, and this is a process error of mine that I want to name precisely rather than wave past.

The one-per-season limit in v1.0 applied to wrongful conviction. Fresh cases were "controlled exceptions" with no number. In round two I *proposed* a one-per-season fresh-case cap as an amendment, which was legitimate. I then, in the same document, listed the consequence of my own unaccepted proposal in the **settled points** table and drew a conclusion from it: that Season One's slot was spent and the pilot could therefore not be a fresh case.

That is a proposal promoted to a premise inside a single document, which is exactly the failure the consensus protocol exists to prevent, and it is the same defect class as an author scoring their own rubric. Withdrawn in full, including the conclusion.

B7. A failed Listener occupies no season slot until it passes new gates.

AGREE, and it is better than my formulation. I had it as a placement; treating a failed story as unallocated until it re-passes is more honest and removes the temptation to schedule around it.

C. Consequence-and-agency ruling

AMEND. Exact final wording below.

First: the objections to my rule are correct, and one of them is fatal on its own

The comparison point was wrong. A case that opens with a person imprisoned, condemned, or denied a body to bury may have **no** credible ending that leaves them below their opening state, and would still be a serious episode. My rule would have failed legitimate stories, and it would have done so most often on exactly the case shapes this backbone is built for. That objection alone is decisive.

The other three hold too. A compliance requirement to produce harm is satisfied most cheaply by a sacrificial branch or a foolish option. Harm as a checkbox produces manufactured harm. And "caused by the player" does read as culpability, when the design goal is responsibility under uncertainty. Conceded, all of it.

Second: the replacement retains the insight and has one gap

Branch-to-branch comparison is the right frame and I adopt it. The counterfactual closing clause is the strongest single sentence in either draft and I keep it verbatim.

The gap is in this clause: *at least one fair branch must realise a material and irreversible cost that another fair branch avoids.*

It does not say **cost to whom, or of what**. A polished retrospective episode can satisfy it with a cost that falls on nobody: a mural is preserved in one branch and lost in the other, irreversibly and materially, while every named person's situation improves in both. The preceding clause about end states differing in who is protected or believed does most of the work, but "differs" is not "loses", and an episode can differ pleasantly.

There is a second gap: nothing requires the cost to be **perceived**. The Listener's branch was mechanically real and testers described it as tacked on, which is what an unfelt consequence feels like. A cost that exists only in the design document is not a cost.

Final wording

Consequence-and-agency gate

Before resolution, at least one material decision must force the player to choose between competing, credible risks affecting named people, or an irreversible evidentiary or publication outcome. **No option may be consequence-free.**

At least two reachable end states must differ materially in who is protected, exposed, believed, harmed, or denied remedy.

At least one fair branch must impose **on a named person the loss of something they held or were owed at the moment of choice**, which another fair branch avoids and which cannot be undone within the episode.

Any harmful outcome must arise causally from a signalled investigative or publication choice, and not from luck, withheld information, or an obviously foolish option.

The player must be able to perceive that loss and attribute it to their own choice before the episode ends. A cost that appears only in the design document is not a cost.

If removing the player's decisions leaves every named person's outcome unchanged, the episode fails.

Two additions to v1.1: the cost is bound to a named person and to a loss rather than a difference, and the loss must be perceptible and attributable. Nothing else changed, and the counterfactual clause is verbatim.

Why this rejects a polished retrospective episode

The episode in section C of my round-two response — every path arrives at some form of *the truth came out*, humane, well made, nothing at risk — now fails on the third clause. Nobody loses anything they held. It cannot be rescued by adding an irreversible-but-victimless difference, because the loss must land on a named person. And it cannot be rescued by putting the loss in the design document, because the player has to feel it.

D. Abuse test

D1. An obviously foolish "publish the witness's address" option creates a bad ending. FAIL, correctly rejected. Excluded by name in clause four. It also fails the fairness test's first question, because no reasonable player would choose it, which means it is not a decision.

D2. Both endings improve a wrongfully imprisoned person's situation, but one loses the only chance of prosecution. PASS. This is the case my original rule wrongly rejected, and the reason the

amendment was necessary. End states differ in who is exposed. The lost prosecution is a loss of something a named person was owed, and the amended wording forces the pitch to say who: the victim's family, named, not "justice" in the abstract. Irreversible, fair, signalled.

D3. A reasonable choice on available evidence, then hidden information produces a death.

FAIL, correctly rejected. Excluded by clause four. This is regret manufactured by withholding, and the fairness test's fourth question exposes it: later evidence may change interpretation, but it may not be the sole cause of the harm.

D4. The player can protect a keeper or preserve decisive evidence, but cannot do both. PASS, and this is the exemplar. Competing credible risks, both affecting named people, no consequence-free option, two end states differing in who is protected and who is believed, an irreversible loss on either branch, and the trade is legible at the moment of choosing. Every clause satisfied without strain. This pattern should be cited in the pitch framework as the reference case.

E. Experiment approval

NOT READY. Four exact changes. Three are unaddressed items from round two; one is a regression introduced in v1.1.

E1. Restore the per-opening observed measure. This is the regression. v1.1 reads all six and then asks. That removes the only behavioural measure in my round-two design and replaces it with recall plus stated choice, which puts opening one at a six-opening disadvantage against opening six on recency alone.

Required change: after each opening, stop and say nothing for five seconds. Record verbatim whether the tester speaks unprompted and what they say. Then proceed. Total cost: thirty seconds across the session. This is the only measure in the design that is observed rather than reported.

E2. Add a blind craft control. The desk gate tests content, not writing quality. Six openings by one author with a preferred outcome will not be equally well written and the difference will not be conscious.

Required change: write nine, not six. A person blind to the hypothesis ranks all nine for writing quality alone, with no knowledge of case family. Discard the three outliers and test the remaining six. If the discard falls unevenly across families, that is itself a finding and must be reported.

E3. Specify the panel. The twelve original testers know the project, know the creator is worried about jeopardy, and will be helpful. Nothing in v1.1 addresses who is in the room.

Required change: at least half the panel must be readers with no history on this project. No tester is told what is being decided until after the reveal step.

E4. Reduce what the opening is required to contain. This is the most likely cause of a null result. v1.1 requires each opening to make legible the live consequence, the rational resistance, the reconstructive activity **and** the first consequential decision, in one hundred and twenty to one hundred and sixty spoken words. Four structural elements in about sixty seconds produces exposition, not drama, and it will produce it in all six openings equally. The test would then measure our ability to write to a checklist and would discriminate nothing, which is the same failure the voided scorecards had.

Required change: the spoken opening must dramatise an event and make **the live consequence** legible. Resistance, reconstruction and the first decision must be answerable at the **desk gate** and need not appear in the spoken text. The desk gate keeps all five of its questions. The opening carries one of them.

What is accepted without change, and two items are improvements on my design: the one hundred and twenty to one hundred and sixty word band, which lands correctly at fifty-five to eighty seconds at restrained pace; randomised order; a single performer; no explanation of family or hypothesis; the forced first, second and stop-on choice; and the two questions "**at what line did you decide?**" and "**who or what did you think was at risk?**", which are better than anything I proposed. The first localises pull to a rewritable moment. The second reveals whether the live consequence landed without ever naming it.

F. Final convergence block

Internal engine: The record is wrong, and somebody is still living with the consequence.

Player-facing promise: Ally Quinn reopens the cases Havenbay calls closed.

(One open item. v1.1 retains "has decided are over." I proposed "calls closed" for the institutional verb and the shorter line. This is a preference, not a rule, it does not gate the experiment, and I will take either. Recorded so it is not lost.)

Episode construction rule: A contested record, a live human consequence, rational resistance, reconstructive play, consequential agency, and an accountable resolution.

Consequence-and-agency gate: as set out in section C above.

Settled. Access improves but is never automatic and evidence custody remains Kill Gate Zero · a passive beneficiary is not an antagonist and every episode needs a rational active resistance system · a deadline is a clock and never resistance · clocks are a separate, optional, case-born field · historical murder plus impunity is a foundation and not present jeopardy · the binary tester question is discarded · the test uses performed openings · reconstruction mode is a required axis · resolution mode is a required season variation field · neither may repeat in consecutive main episodes · Slack Water stays outside the argument until its pitch is in the shared packet · The Listener remains demoted and holds no season slot until it re-passes · no fresh-case quota was ever proposed and my derived conclusion from one is withdrawn · the test decides dominance, marketing promise and pilot family, not permission, and exclusivity is not recommended by anyone.

Unresolved. One: whether section C's amended wording is accepted. Everything else in this process is closed.

(Not unresolved but pending: the four experiment changes in section E, which are execution details rather than disagreements, and the promise-wording preference above.)

Next action. Accept or reject the section C wording. If accepted, apply the four section E changes, then write nine openings to the template, blind-rank them, and test six.

PART TWO — THE CONSENSUS DOCUMENT

Everything below is the merged state of the process. It is written to be liftable into canon once the section C wording is accepted and the spoken-opening test has run. **It is not canon yet and no project file has been changed.**

1. The three levels

Level	Wording	Used by
Internal engine	The record is wrong, and somebody is still living with the consequence.	Writers, at premise stage
Player-facing promise	Ally Quinn reopens the cases Havenbay calls closed. (<i>v1.1 alternative: "has decided are over". Preference open, not gating.</i>)	Store copy, marketing, first-run framing
Episode construction rule	A contested record, a live human consequence, rational resistance, reconstructive play, consequential agency, and an accountable resolution.	Production gate, at pitch stage

2. The six mandatory elements

Every main episode must contain all six. Any one missing is a REWORK.

1. **Contested record.** An official, accepted or socially dominant account that can be stated plainly in one sentence.
2. **Live human consequence.** A named living person under a consequence that is still operating. Not a category. Not historical suffering that has stopped.
3. **Rational resistance.** An agent or interest that benefits from, depends on, or is protected by the accepted record, and that responds to being investigated. Never a deadline. Its restraint must be as rational as its action.
4. **Reconstructive play.** The merge requirements *are* the reconstruction, not the price of reaching it. The physical activity is named at pitch stage.
5. **Consequential agency.** The section 4 gate, in full.
6. **Accountable resolution.** The dramatic question is answered and publication is reckoned with. A conviction is not required.

3. Access map, mandatory pitch fields

For every decisive item of evidence:

1. Who lawfully possesses it.
2. Why that person or body can disclose it.
3. What remains inaccessible.
4. How Ally proceeds without police privilege.
5. **What Ally gets wrong because she can only reach half the record.**

Field five is not compliance. The inaccessible half is where the episode's wrong hypothesis lives, and every audit this project has run has marked every pitch down for not having one.

4. Consequence-and-agency gate

Before resolution, at least one material decision must force the player to choose between competing, credible risks affecting named people, or an irreversible evidentiary or publication outcome. **No option may be consequence-free.**

At least two reachable end states must differ materially in who is protected, exposed, believed, harmed, or denied remedy.

At least one fair branch must impose **on a named person the loss of something they held or were owed at the moment of choice**, which another fair branch avoids and which cannot be undone within the episode.

Any harmful outcome must arise causally from a signalled investigative or publication choice, and not from luck, withheld information, or an obviously foolish option.

The player must be able to perceive that loss and attribute it to their own choice before the episode ends. A cost that appears only in the design document is not a cost.

If removing the player's decisions leaves every named person's outcome unchanged, the episode fails.

Reference case, to be cited in the pitch framework: the player can protect a keeper or preserve decisive evidence, and cannot do both.

Fairness test, per consequential decision

The pitch must state, for each: what the player knows at the moment of choice · which risk each option carries · why a reasonable player might choose either · what later evidence changes the interpretation · which consequence cannot be undone.

If one option is obviously correct before the choice, it is not agency. If the decisive information is withheld solely to manufacture regret, it is not fair play.

Jeopardy gate, retained from round two

Both tracks required, judged on the weaker. **Human consequence:** a named living person, a mechanism still running, stated. **Investigative consequence:** a specific, stated thing the player's choices can make worse. Physical danger is one valid form of either and is neither required nor the default. The human consequence must be stateable in one sentence at listening speed, present in the cold open for the pilot and no later than the first podcast milestone otherwise.

5. The three-axis taxonomy

A pitch names one value on each axis. **Neither axis two nor axis three may repeat in consecutive main episodes.**

Axis one, what is wrong with the record: never solved · wrongful conviction · misclassified · covered up · officially solved but contested. *Wrongful conviction is provisionally limited to one main episode per season, for formula risk, for research burden, and because it pulls Ally toward a lawyer's investigator when podcaster leads the player-facing identity.*

Axis two, what the player reconstructs: a recording · a document set · an environmental model such as tide, weather, visibility or travel time · a person's identity · a timeline · a route or voyage · a place that no longer exists · the custody history of an object · a network of relationships or payments.

Axis three, how the episode resolves: legal review or exoneration · protection of a keeper or future victim · public correction without prosecution · institutional exposure · confession or confrontation · recovery, identification or return · a deliberately limited publication to prevent further harm · a partial remedy with a named cost · a culprit exposed but not yet chargeable.

Live consequence, paired with axis one, not itself an anti-formula axis: wrongful loss of liberty or status · continued access or power · keeper under pressure · evidence opportunity closing · institutional decision · public lie blocking remedy.

6. Fresh cases

Permitted, not preferred. No numeric quota. A fresh case must satisfy the same six elements, the same access map, the same jeopardy gate and the same agency gate. The season mix is an outcome of the spoken-opening test and later full-pitch ranking, not a rule set in advance.

7. Clocks

Separate, optional, and only where the case supplies one: a scheduled hearing, transfer, lawful disposal, demolition, appointment, payment or official decision. Pause-safe per the standing clock convention. Not required where current harm is already accumulating. Never counted as resistance.

8. The spoken-opening experiment

Purpose. Decides which family has the strongest pilot pull, and informs whether contested records dominate the main-episode slate. Does **not** decide whether contested records, reconstruction, lawful access maps or pause-safe clocks may be used, all of which are approved now. Does not impose exclusivity, which nobody recommends.

Materials. Write **nine** openings: three fresh murders, three never-solved or misclassified, three officially solved but contested. Each one hundred and twenty to one hundred and sixty spoken words, VO-lock compliant, parsing at listening speed, no titles, no category labels.

Each opening **dramatises an event** and makes **the live consequence** legible. It does not have to carry resistance, reconstruction or the first decision in its spoken text; those are answered at the desk gate.

Desk gate. Reject before table read if the pitch behind an opening cannot answer: what happens in the first thirty seconds · who can be materially affected now · what rational force resists correction · what the player's hands will reconstruct · what first choice changes the case rather than a line of dialogue.

Blind craft control. A person blind to the hypothesis ranks all nine for writing quality alone, without knowing case families. Discard the three lowest. If the discards fall unevenly across families, report it as a finding.

Panel. At least five, ideally more. **At least half must have no history with this project.** Nobody is told what is being decided until the reveal step.

Method. Read all six aloud, one performer, restrained delivery, randomised order per tester. **After each opening, stop and say nothing for five seconds; record verbatim whether the tester speaks unprompted and what they say.** Then, after all six: force the one they would play first, the one they would try second, and the one they would stop on. Then ask, in order: *what happened that made you choose that one* · *at what line did you decide* · *who or what did you think was at risk*. Then, last, reveal the backbone sentence and ask whether it makes the series more or less attractive.

Never ask whether they prefer cold cases or fresh murders.

Decision rule, no scores, no ties. Each tester names exactly one first, one second, one stop-on. A refusal is recorded as a refusal, not a tie; refusals above a third of the panel mean the openings are not discriminating and must be rewritten before any conclusion.

- **Contested records dominate** if they take first choice from a clear majority and are at least equal on unprompted speech, with stated reasons tied to injustice, consequence, discovery or reconstruction.
- **Mixed season, no dominance** if fresh openings take the forced choice while contested openings match or beat them on unprompted speech. Both families pull; the difference is delivery, not substance.
- **Backbone rejected as the dominant slate** if fresh openings win the forced choice *and* unprompted speech *and* the stated reasons cluster on immediacy and danger. All three.
- **Copy rejected, not the backbone** if contested openings become attractive only after the backbone sentence is revealed.
- **Desk failure, no test needed** if either contested pair cannot be written to the fresh pair's standard. The dominance hypothesis fails before anyone is in the room.

9. Document impact, if adopted after the test

Sequenced, pivot brief first, because every other document defers to it: pivot brief (high concept, pilot doctrine, jeopardy, clock examples, case-family language) · critical reviewer charter (matching gates) · story-generation prompt to v2.1 (evidence custody fields, three-axis taxonomy, agency gate) · bible (high concept, Ally's practice, pilot doctrine, season architecture, superseded twelve-lead material) · bible §6.4 terminology collision with between-episode "cold cases" · season straw man · pitch framework (access map, resistance, clock, agency, axes as mandatory fields) · marketing promise.

No canon file changes until the test runs or the creator explicitly accepts the remaining risk.

10. Status

Status	Items
Settled	All fourteen items listed in Part One section F.
Open, one item	Whether Part One section C's amended agency wording is accepted.
Pending, not disputed	The four experiment changes in Part One section E; the player-promise wording preference; the Slack Water pitch entering the shared packet.

Revision history

Version	Date	Change
1.0	2026-08-24	First cross-model consensus brief. Two-axis taxonomy, revised jeopardy gate, blind shelf test.
1.1	2026-08-24	Incorporated round-two withdrawals; added player-reconstruction and resolution variation; separated resistance from clocks; corrected the fresh-case quota misread; proposed the consequence-and-agency gate.
1.2	2026-08-24	Merged document. Adjudicates v1.1, adopts the branch-to-branch agency frame with two amendments (loss bound to a named person; perceptibility), withdraws the fresh-case cap and its derived conclusion, withdraws "exclusivity", adopts resolution mode as a third variation axis, and marks the experiment NOT READY with four required changes including restoration of the per-opening observed measure.