

ALLY QUINN STORY BACKBONE —

CONSENSUS v1.3 (FINAL)

Document ID: AQ-NAR-BACKBONE-CONSENSUS · **Version:** 1.3 · **Date:** 2026-08-24 · **Status:** CONSENSUS REACHED · **Supersedes:** v1.0, v1.1, v1.2 and the round-two adjudication

Canon impact: none yet. No project file was edited in any round of this process. This document is written to be lifted into canon after the spoken-opening test, or earlier if Stephen accepts the remaining risk.

o. Consensus statement

Both reviewers accept every clause below. **Nothing in this document is disputed.**

The two final amendments were made by the second reviewer and are adopted verbatim:

1. **"Material loss", bound to the end-state difference.** My wording allowed a named person to lose something trivial while both endings remained safe. The loss must now be material and must itself constitute one of the differences between reachable end states.
2. **Retain the top two per family rather than discard the lowest three.** My blind-craft control was correct and my selection rule introduced a fresh confound: discarding the three lowest regardless of family can leave a slate of three, two and one, and a family with more openings on the table has more chances to be chosen in a forced pick, independent of quality.

The player-facing promise is settled as **"has decided are over"** rather than my "calls closed", on the second reviewer's reasoning, which is better than mine: the broader phrase covers formally closed cases, abandoned investigations, socially accepted explanations and stories the city prefers not to revisit. "Closed" implies a formal file status and would exclude two of the five record shapes in the taxonomy this document adopts.

One implementation note, not a change and not a dispute. The retention rule permits a desk failure where "a family cannot supply two openings that meet the common craft standard." A ranking always produces a top two, so as written that branch can never trigger. The blind reviewer should therefore mark each of the nine **keep or do not keep** in addition to ranking them, which supplies a threshold independent of position. This does not alter the rule, and I flag it only so whoever runs the ranking is not left to invent a standard on the day.

1. The three levels

Level	Wording	Used by
Internal engine	The record is wrong, and somebody is still living with the consequence.	Writers, at premise stage
Player-facing promise	Ally Quinn reopens the cases Havenbay has decided are over.	Store copy, marketing, first-run framing
Episode construction rule	A contested record, a live human consequence, rational resistance, reconstructive play, consequential agency, and an accountable resolution.	Production gate, at pitch stage

2. The six mandatory elements

Every main episode must contain all six. Any one missing is a REWORK.

1. **Contested record.** An official, accepted or socially dominant account, stateable plainly in one sentence.
2. **Live human consequence.** A named living person under a consequence that is still operating. Not a category. Not historical suffering that has stopped.
3. **Rational resistance.** An agent or interest that benefits from, depends on, or is protected by the accepted record, and that responds to being investigated. Never a deadline. Its restraint must be as rational as its action.
4. **Reconstructive play.** The merge requirements *are* the reconstruction, not the price of reaching it. The physical activity is named at pitch stage.
5. **Consequential agency.** The section 4 gate, in full.
6. **Accountable resolution.** The dramatic question is answered and publication is reckoned with. A conviction is not required.

3. Access map, mandatory pitch fields

For every decisive item of evidence:

1. Who lawfully possesses it.
2. Why that person or body can disclose it.
3. What remains inaccessible.
4. How Ally proceeds without police privilege.
5. **What Ally gets wrong because she can only reach half the record.**

Field five is not compliance. The inaccessible half is where the episode's wrong hypothesis lives, and every audit this project has run has marked every pitch down for not having one.

4. Consequence-and-agency gate

Before resolution, at least one material decision must force the player to choose between competing, credible risks affecting named people, or an irreversible evidentiary or publication outcome. **No option may be consequence-free.**

At least two reachable end states must differ materially in who is protected, exposed, believed, harmed, or denied remedy.

At least one fair branch must impose on a named person a **material loss** of something they held or were owed at the moment of choice. **That loss must constitute one of the material differences between the reachable end states**, another fair branch must avoid it, and it cannot be undone within the episode.

Any harmful outcome must arise causally from a signalled investigative or publication choice, and not from luck, withheld information, or an obviously foolish option.

The player must be able to perceive that loss and attribute it to their own choice before the episode ends. A cost that appears only in the design document is not a cost.

If removing the player's decisions leaves every named person's outcome unchanged, the episode fails.

Reference case, to be cited in the pitch framework: the player can protect a keeper or preserve decisive evidence, and cannot do both.

Verified against four abuse patterns, both reviewers agreeing

Pattern	Result	Why
An obviously foolish "publish the witness's address" option produces a bad ending	FAIL, correctly rejected	Excluded by name. No reasonable player chooses it, so it is not a decision
Both endings improve a wrongfully imprisoned person, but one loses the only chance of prosecution	PASS	The victim's named family loses a material remedy they were owed; the loss distinguishes the endings and another fair branch avoids it
A reasonable choice on available evidence, then hidden information produces a death	FAIL, correctly rejected	Harm from concealment rather than from a signalled risk
Protect the keeper or preserve decisive evidence, not both	PASS, the exemplar	Competing credible risks, no consequence-free option, legible trade, irreversible either way

Fairness test, per consequential decision

The pitch must state: what the player knows at the moment of choice · which risk each option carries · why a reasonable player might choose either · what later evidence changes the interpretation · which consequence cannot be undone.

If one option is obviously correct before the choice, it is not agency. If the decisive information is withheld solely to manufacture regret, it is not fair play.

Jeopardy gate

Both tracks required, judged on the weaker.

Human consequence. A named living person is under a consequence still operating: confined, excluded, publicly condemned, exposed, unprotected, denied a remedy, financially trapped, or reachable by someone with reason to reach them. State the person, the mechanism, and why it is still running. Historical suffering that has stopped fails. A category rather than a person fails.

Investigative consequence. The player's choices can make a specific, stated thing worse: evidence lost or made inadmissible · a witness permanently closed to the show · a hearing or prosecution prejudiced · a person exposed before they were ready · a remedy delayed past its opportunity · Ally's ability to be trusted, to publish, or to work lawfully permanently reduced. Embarrassment without consequence fails. Audience numbers alone fail. Reversible within the episode fails.

Legibility. The human consequence must be stateable in one sentence at listening speed: present in the cold open for the pilot, and no later than the first podcast milestone otherwise.

Physical danger is one valid form of either track. It is not required and it is not the default.

5. The three-axis taxonomy

A pitch names one value on each axis. **Neither axis two nor axis three may repeat in consecutive main episodes.**

Axis one, what is wrong with the record. Never solved · wrongful conviction · misclassified · covered up · officially solved but contested. *Wrongful conviction is provisionally limited to one main episode per season: formula risk, research burden, and because it pulls Ally toward a lawyer's investigator when podcaster leads the player-facing identity.*

Axis two, what the player reconstructs. A recording · a document set · an environmental model such as tide, weather, visibility or travel time · a person's identity · a timeline · a route or voyage · a place that no longer exists · the custody history of an object · a network of relationships or payments.

Axis three, how the episode resolves. Legal review or exoneration · protection of a keeper or future victim · public correction without prosecution · institutional exposure · confession or confrontation · recovery, identification or return · a deliberately limited publication to prevent further harm · a partial remedy with a named cost · a culprit exposed but not yet chargeable.

Live consequence, paired with axis one and not itself an anti-formula axis. Wrongful loss of liberty or status · continued access or power · keeper under pressure · evidence opportunity closing · institutional decision · public lie blocking remedy.

6. Fresh cases

Permitted, not preferred. **No numeric quota.** A fresh case must satisfy the same six elements, access map, jeopardy gate and agency gate. The season mix is an outcome of the spoken-opening test and later full-pitch ranking, never a rule set in advance.

7. Clocks

Separate, optional, case-born: a scheduled hearing, transfer, lawful disposal, demolition, appointment, payment or official decision. Pause-safe per the standing clock convention. Not required where current harm is already accumulating. **Never counted as resistance.**

8. The spoken-opening experiment

Scope. Decides which family has the strongest pilot pull, and informs whether contested records dominate the main-episode slate. Does **not** decide whether contested records, reconstruction, lawful access maps or pause-safe clocks may be used: all are approved now. Does not impose exclusivity, which nobody recommends.

Materials. Write **nine** openings: three fresh murders, three never-solved or misclassified, three officially solved but contested. Each one hundred and twenty to one hundred and sixty spoken words, VO-lock compliant, parsing at listening speed, no titles, no category labels.

Each opening **dramatises an event** and makes **the live consequence** legible. It need not carry resistance, reconstruction or the first decision in its spoken text; those are answered at the desk gate.

Desk gate. Reject before table read if the premise behind an opening cannot answer: what happens in the first thirty seconds · who can be materially affected now · what rational force resists correction · what the player's hands will reconstruct · what first choice changes the case rather than a line of dialogue.

Blind craft control and slate balance. A person blind to the hypothesis ranks all nine for writing quality alone, without knowing case families, and marks each **keep or do not keep**. Retain the **two highest-ranked from each family**. If a family cannot supply two marked keep, rewrite that family or record a desk failure. **Do not test an unbalanced slate.** Uneven ranking is a finding and is reported, but is never a reason to run a confounded test.

Panel. At least five, ideally more. **At least half must have no history with this project.** Nobody is told the hypothesis until every behavioural and forced-choice response has been recorded.

Method. Read the six aloud, one performer, restrained delivery, randomised order per tester. **After each opening, stop and say nothing for five seconds. Record verbatim whether the tester speaks unprompted and what they say.** This is the only observed measure in the design and it is never converted into a score.

Then, after all six: force the one they would play first, the one they would try second, and the one they would stop on. Then ask, in order: *what happened that made you choose that one · at what line did you decide · who or what did you think was at risk.*

Then, last, reveal the backbone sentence and ask whether it makes the series more or less attractive.

Never ask whether they prefer cold cases or fresh murders.

Decision rule. No scores, no percentages, no ties. Each tester names exactly one first, one second, one stop-on. A refusal is recorded as a refusal, not a tie. Refusals above a third of the panel mean the openings are not discriminating and must be rewritten before any conclusion is drawn.

- **Contested records dominate** if they take first choice from a clear majority and are at least equal on unprompted speech, with reasons tied to injustice, consequence, discovery or reconstruction.

- **Mixed season, no dominance** if fresh openings take the forced choice while contested openings match or beat them on unprompted speech. Both pull; the difference is delivery, not substance.
- **Backbone rejected as the dominant slate** if fresh openings win the forced choice **and** unprompted speech **and** the reasons cluster on immediacy and danger. All three.
- **Copy rejected, not the backbone** if contested openings become attractive only after the backbone sentence is revealed.
- **Desk failure, no test needed** if a contested family cannot be written to the fresh family's standard.

9. Consequence for the pitch slate

All six competing pitches must be generated or conformed under the same rules. The three existing challengers and the surviving incumbent were produced under prompt kit v2.0, which predates everything in sections 2 to 5. Comparing them against pitches built to v2.1 would weigh new machinery against an old gearbox.

Prompt kit v2.1 must therefore be issued before any further pitch generation, incorporating: the consequence-and-agency gate · the three-axis taxonomy with the consecutive-repeat rule · the five-field access map · separate resistance and clock fields · the reconstruction requirement that merge requirements *are* the reconstruction · resolution-mode variation.

A sequencing consequence both reviewers passed over. The nine openings cannot be written from nothing: each needs a premise capable of answering the desk gate's five questions. Only about three such premises currently exist. Nine premise sketches are therefore required first, and they should be generated under v2.1 so they pass the desk gate by construction rather than by repair. **v2.1 comes before the openings, not after them.**

10. Document impact, when adopted

Sequenced, pivot brief first, because every other document defers to it: pivot brief (high concept, pilot doctrine, jeopardy, clock examples, case-family language) · critical reviewer charter (matching gates) · prompt kit v2.1 · bible (high concept, Ally's practice, pilot doctrine, season architecture, superseded twelve-lead material) · bible §6.4 terminology collision with between-episode "cold cases" · season straw man · pitch framework (access map, resistance, clock, agency, three axes as mandatory fields) · marketing promise.

No canon file changes until the test runs or the creator explicitly accepts the remaining risk.

11. What each reviewer withdrew

Recorded because the protocol asks for explicit convergence rather than mutual politeness, and because knowing what was wrong is worth more later than knowing what was agreed.

First reviewer withdrew: that a closed case is public · that the antagonist need not act · that the clock problem ends · that a free killer is permanent jeopardy · the binary tester question · the worse-than-open rule · the one-fresh-case cap and the conclusion derived from it inside the same document · the word "exclusivity" · the claim that reconstruction is the only directly experienced axis · the discard-lowest-three selection rule.

Second reviewer amended, and was upheld on all of it: the two-axis taxonomy needed a third and then a fourth field · the resistance list wrongly included deadlines · the fresh-case quota was a misreading · the comparison point in the agency rule was unsafe · the loss required binding to the end-state difference · the slate required balancing.

Neither reviewer's position survived intact. That is the intended outcome of the protocol.

Revision history

Version	Date	Change
1.0	2026-08-24	First cross-model consensus brief. Two-axis taxonomy, revised jeopardy gate, blind shelf test.
1.1	2026-08-24	Incorporated round-two withdrawals; player-reconstruction and resolution variation; resistance separated from clocks; fresh-case quota misread corrected; consequence-and-agency gate proposed.
1.2	2026-08-24	Merged document. Branch-to-branch frame adopted with two amendments; fresh-case cap and derived conclusion withdrawn; experiment marked NOT READY with four required changes.
1.3	2026-08-24	CONSENSUS REACHED. "Material loss" bound to the end-state difference; slate balanced by retaining the top two per family; player promise settled as "has decided are over"; keep-or-not threshold added as an implementation note; v2.1-before-openings sequencing recorded.