

MODE A CROSS-AUDIT INSTRUMENT — CAI-1

FINAL

Instrument ID: CAI-1 · **Status:** FROZEN on contestation. Amendments of 2026-08-25 applied in full · **Date:** 2026-08-25

Purpose. A single shared instrument for auditing two existing Mode A executions against the same evidentiary standard in both directions. It is a checklist, not a controlled prompt, and it creates no premises.

Subjects. `mode-a-EX-2026-08-24-02` (nine premises, Claude) and `AQ-NAR-MODEA-EXEC-2026-08-24-01` (nine premises, ChatGPT). Both frozen as evidence. Neither may be rewritten.

Direction. Both auditors conduct a **complete item-by-item audit** of the other's nine. New research may test, contradict or verify support already supplied by the execution. **It may not supply missing support, repair a premise, or convert NOT DEMONSTRATED into CONFORMS.** A foundational claim left unresolved in the execution remains subject to A4.

Freeze protocol. Contest this draft. When both parties settle the wording, it freezes as **CAI-1 FINAL** and neither party may amend it mid-audit. A defect discovered in the instrument during the audit is recorded and applied in a successor instrument, not retrofitted.

1. Outcomes

Every item receives exactly one.

Outcome	Meaning
CONFORMS	The execution record itself demonstrates the requirement. Not "could be made to."
FAILS	A verified contradiction, an unknown foundational mechanism, or a gate substantively unmet.
NOT DEMONSTRATED	The premise may be viable; the record does not contain enough evidence to award CONFORMS.

OUT OF INSTRUMENT is not an item outcome. It is a separate status for concerns falling outside items A1 to I7. Recorded, never determinative, carried to successor design.

Conditional items. For B5 and I1, where the trigger is optional: **explicitly demonstrated non-use CONFORMS.** If the record does not establish whether the trigger occurred, the outcome is NOT DEMONSTRATED.

Decision rules, binding on both auditors.

- No repair during audit.** An auditor may not supply a missing answer, however obvious.
- NOT DEMONSTRATED never becomes CONFORMS through charitable inference,** and never becomes FAILS through uncharitable inference. It is its own finding.

3. **An auditor may not substitute preference for evidence.** Where the record answers a requirement in a way the auditor would not have chosen, the outcome is CONFORMS. Disagreements of craft are OUT OF INSTRUMENT.
 4. **Every finding cites a quotation or a precise record location.** Where a required element is absent, cite the section searched and identify the absent element. **Missing evidence is NOT DEMONSTRATED** unless the requirement is explicitly an output-completeness requirement, or the record confirms a substantive failure.
 5. **No scores, no percentages, no ranking.**
 6. **The auditor does not decide advancement.** Auditors report item outcomes and a premise-level conformance summary. Advancement is the owner's call at adjudication. This keeps selection out of the auditor's hands and prevents quota-shaped outcomes re-entering through the audit.
-

2. Per-premise items

A. Kill Gate Zero

#	Requirement
A1	A first-five- causal -event audit is present, five events, not five scenes
A2	All eight fields present for every event: actor, authority, custody, normal response, safeguard, bypass, consequence, evidence source
A3	No field blank, circular or invented
A4	No premise-dependent authority, custody, access, procedural, technological or institutional claim is left UNKNOWN. A deferral to "verify before script" is a FAIL under this item, not a mitigation
A5	All institutional-power rules respected: corruption can misuse an existing power but cannot create one · a contractor may perform a delegated physical task only after valid authority is granted, and a contract does not automatically displace police, court, coroner, medical or regulatory control · a forged document may deceive a plausible recipient for a plausible period but does not automatically neutralise independent checks, custody records, signatures, witnesses or system alerts · a database entry can mislead action but cannot make a living person legally or physically inaccessible without a credible downstream mechanism · police, coroners, doctors, forensic staff, journalists, archivists and digital specialists remain competent and their competence creates constraints the antagonist must overcome · Ally's PI licence grants no police powers, no unrestricted crime-scene access, no medical records, no protected digital systems and no custody of evidence
A6	Antagonist least-cost test answered: why murder, why now, why this victim, why this method and staging, why not intimidation, theft, flight, bribery, destruction or silence by an easier route, and what the method costs the antagonist
A7	The six commonsense questions are answered in one plain sentence each, not merely answerable: why do the police not prevent this · why does the coroner, doctor, court, employer, bank, platform or regulator permit this · why is Ally present and allowed to do what the story requires · why does the antagonist choose this method instead of a simpler and safer one · why does the victim or keeper not take the obvious protective action · what would an informed viewer object to in the first story beat
A8	Every premise-dependent operational claim is supported near the claim by an authoritative and preferably primary source. News, blogs, entertainment

#	Requirement
	references and general summaries may locate primary material but may not provide final support. Jurisdictional differences and their materiality are stated

B. Gates one to nine

#	Requirement
B1	Gate 1 desk: planned first thirty seconds are murder-led, recognisably true crime, contain a concrete contradiction, danger or irreversible event, and survive the ordinary-viewer objection. Performed compulsion recorded UNKNOWN
B2	Gate 2 human consequence: a named living person , a mechanism still operating , both stated. A category fails. Historical suffering that has stopped fails
B3	Gate 2 investigative consequence: a specific stated thing the player's choices can make worse. Embarrassment without consequence fails. Audience numbers alone fail. Reversible-within-episode fails
B4	Gate 2 legibility: the human consequence is stateable in one sentence at listening speed
B5	Gate 3 clock: optional. If used, case-born, pause-safe, not a session-contained window, and not the same object as the resistance
B6	Gate 4 midpoint: the player's action, route, target, risk or evidence strategy changes, not only their interpretation
B7	Gate 5 fair play: conclusions traceable to obtained evidence; a fair early plant precedes its explanation; no coincidence proves the culprit; a plausible wrong hypothesis is present
B8	Gate 6 humanisation: the victim is grievable through one concrete human detail before becoming a puzzle component ; the keeper has a need beyond survival or delivering information; the criminal economy imposes a recognisable ordinary cost ; no gore, child harm or melodrama
B9	Gate 7, v2.1 verbatim: could this exact case happen only in Havenbay? Atmosphere must advance evidence, access, motive, jeopardy or consequence. Fog, docks, old records and civic corruption are not sufficient by themselves
B10	Gate 8, in full: fifteen to twenty played leads across two or three phases determined by the story · played and total authored leads stated separately · Field Kit use, any Lab unlock and Kestrel Corner function identified · approved, proposed and new portraits and locations identified · reconstruction and resolution modes named · no invented costs or economy figures

#	Requirement
B11	Gate 9 resistance: an agent or interest that benefits from, depends on or is protected by the accepted record and responds when challenged . Not a deadline, schedule or passage of time. Not a passive evidentiary constraint. Restraint as rational as action

C. Access map

#	Requirement
C1	All five fields present for the decisive item
C2	A lawful possessor is named
C3	A disclosure route is stated and is plausible on the supplied research
C4	What remains inaccessible is stated
C5	Ally's route requires no police privilege and no power she does not have
C6	Field five present: what Ally gets wrong because she can reach only part of the record

D. Agency

#	Requirement
D1	At least four decisions that change which leads spawn, what the player does next, what evidence survives, the risk state, or the ending. Dialogue-only variation does not count
D2	At least two are Class A
D3	No more than one is Class C
D4	Every branch rejoins within three leads; no fork opens before its branch rejoins
D5	At least two reachable end states differ materially in who is protected, exposed, believed, harmed or denied remedy
D6	At least one fair branch imposes on a named person a material loss of something held or owed at the moment of choice; that loss is one of the end-state differences; another fair branch avoids it; it cannot be undone within the episode
D7	The player can perceive that loss and attribute it to their own choice before the episode ends
D8	Counterfactual: removing the player's decisions changes at least one named person's outcome
D9	No consequence-free option; no obviously correct option; no harm arising from luck, withheld information or an obviously foolish option
D10	Fairness test present per consequential decision: what the player knows, the risk each option carries, why a reasonable player might choose either, what later evidence reframes, which consequence cannot be undone

Auditor note on D9. A protect-now versus preserve-evidence choice is **NOT DEMONSTRATED** unless the record establishes competing credible risks and explains why a reasonable player might choose either. **The opposing risk may affect another named person or may constitute an irreversible evidentiary or publication outcome. A second named person is not mandatory**, since the v2.1 reference pattern itself sets a keeper against decisive evidence. Both executions used variants of this shape; it is flagged so it is judged identically in both directions.

E. Registers

#	Requirement
E1	Failure mechanism named from the eight below
E2	Inherited situation named from the seven below
E3	Primary payoff named from the seven below, classified by the counterfactual test below
E4	Axis one, axis two and axis three values named
E5	The merge requirements are the reconstruction. The record states what the player physically does, what conclusion that activity makes possible, and what action the conclusion unlocks. Asserting that the requirements "are the investigation" without stating the physical activity does not conform

The eight failure mechanisms — why the record is wrong. 1. Correct system, wrong answer. Every actor performs correctly and the outcome is still false. 2. Misuse of a power actually held. A discretion, access or authority the person genuinely has, exercised for a reason. 3. Overruled by someone with standing. The system worked and lawful authority set it aside. 4. Unequal looking. Nobody looked hard because of who the victim was. 5. A plausible person lied to the system, and the system believed them. 6. A handoff gap. Two institutions each did their job and the space between them swallowed it. 7. A community enforced the account. Everyone knew, nobody said. 8. The victim's own choice created the record.

The seven inherited situations — what condition does the player enter? 1. Wrong person blamed or punished. 2. Culprit or beneficiary protected by the accepted account. 3. Person, identity, body or decisive evidence missing. 4. Preventable harm still occurring. 5. Keeper, witness or evidence opportunity under pressure. 6. Institution continuing an ordinary harmful practice. 7. Responsibility genuinely unresolved rather than merely concealed.

The seven primary payoffs — what does the player feel they accomplished? 1. Cleared or restored someone. 2. Caught or proved someone responsible. 3. Found or named someone or something. 4. Prevented a further harm. 5. Protected a person, source or evidentiary opportunity. 6. Exposed an institution or criminal economy. 7. Understood responsibility or decided what should become public.

The counterfactual classification test, binding on E3. Classify by effect, not by pitch language. **If removing the blamed person removes the story's main emotional reason to continue, the payoff is clearing or restoring, even if the episode also identifies a killer.**

F. Territory and canon

#	Requirement
F1	The premise does not duplicate or reskin an occupied candidate's protected spine or combination of devices . Individual evidence types remain available unless v2.1 expressly excludes them
F2	Not the excluded structural shape: a damaged or interrupted recording reconstructed by audio analysis to locate a person already missing
F3	Canon characters used consistently with their established functions and status. Portraits and locations accurately labelled approved, proposed or new. No invention left unflagged
F4	Endearment map respected; no new endearment users; Vega's pronouns unset unless flagged as a proposed ruling
F5	Season notch: one new fact and one new cost; no unmasking; no casual connection of the episode killer to the season conspiracy; the coin is not a substitute for an economy
F6	Episode-specific people, places and institutions labelled PROPOSED CANON
F7	Where a premise uses a recording, acquisition, disclosure, police or legal use, consent and publication are assessed separately under Havenbay's all-party-consent rule. Lawful creation or possession does not imply lawful broadcast. The Tip Line's caller-consent rule is preserved

G. Kill brief

#	Requirement
G1	All four objection frames present: competent professional, rational antagonist, ordinary player, producer
G2	Ends PASS or FAIL with a single decisive reason. No conditional pass
G3	No decisive weakness remains unresolved when the brief ends PASS. Where a brief identifies motive, resistance, access or proof as thin, weak, unverified or admitted and then passes without documenting the single permitted core-preserving revision and re-check , G3 FAILS. A weakness repaired once and re-audited may conform

G3 exists because it is the failure this cross-audit was created by. Candour is a finding, not a dispensation. It preserves v2.1's single permitted revision, so a premise genuinely repaired once is not punished for having been candid about the repair.

3. Field-level items

#	Requirement
H1	Seed-level distributions for mechanism, inherited situation and primary payoff. The auditor classifies every retrievable frozen seed and survivor under the three registers; classification is measurement, not premise repair. Seeds not present in the frozen record may not be reconstructed from memory
H2	Survivor-level distributions reported for the same three
H3	The delta between them reported, with the gate responsible for each category's attrition attributed
H4	No more than half the survivor pool shares one inherited situation or one primary payoff
H5	Seed level, per family: at least four inherited situations and four primary payoffs present, and no more than three of eight sharing one

H3 is the instrument's only diagnostic item. Where seeds are diverse and survivors are not, the gates are performing the selection, and that is the finding both executions exist to expose.

4. Execution-level items

#	Requirement
I1	If a second field was invoked, it contained twelve seeds distributed across the deficient families
I2	A complete seed ledger is present, with the single real-world process dependency stated for every seed
I3	The rotating interleaved order was followed and reported: rounds of three, three and two per family, starting family rotating, only the stated batch generated at each turn
I4	The five closest rejected seeds are given with a decisive reason each, unbeautified
I5	true-crime-genre-anatomy not used, and the execution record says so
I6	Mode A prohibitions observed: no cold opens, titles, elevator pitches, rankings, incumbent comparisons or table-read recommendations
I7	Execution record complete against the v2.1 field list

5. Declared conditions before the audit begins

Claude's execution cannot satisfy I2. Its thirty seeds were not recorded as an artifact. Fourteen are retrievable from the published document: nine survivors and five named rejections. The remaining sixteen exist only in unretrievable reasoning. They will not be reconstructed, classified retrospectively or improved.

Pre-assigned outcomes for that execution, agreed before the audit begins. UNAVAILABLE is the *reason*, never a fifth outcome:

Item	Outcome	Reason
H1	NOT DEMONSTRATED	Source material unavailable
H3	NOT DEMONSTRATED	Delta uncomputable without seed-level data
H5	NOT DEMONSTRATED	Seed-level distribution unavailable
I2	FAILS	Output-completeness requirement; the ledger was required and is absent

ChatGPT's execution contains its complete twenty-four-seed ledger with dependencies, so its seed-to-survivor delta is calculable and H1, H3 and I2 are assessable.

Asymmetry recorded: the delta instrument can only run in one direction on the existing evidence. That is a limitation of Claude's execution, not of the instrument, and no inference about Claude's gates-versus-generation attractor may be drawn from its absence.

6. Reporting format

Per premise: a table of every item A1 to G3 with its outcome and, for anything other than CONFORMS, a quotation from the record and one sentence of reason.

Per execution: the field-level items H1 to H5, the execution-level items I1 to I7, and a **premise-level conformance summary** — not an advancement recommendation.

Auditors exchange complete results. Only actual disagreements are adjudicated. Agreement is not restated.

7. What this instrument does not do

It does not rank, recommend, allocate or select. It does not decide which premises advance. It does not assess performed cold-open compulsion, which remains UNKNOWN for every premise in both executions until blind testing.

And it does not extend itself. A concern outside items A1 to I7 is recorded OUT OF INSTRUMENT and carried forward. It does not become a new item mid-audit, and it does not postpone advancement on its own.