

CRITICAL REVIEWER ON STRUCTURE v2.0: REWORK (2026-09-01)

The scheduled Truth-Seeking Critical Reviewer cloud run (trig_01F7v7aCNF773DPj9p3H3GAq) fired 2026-09-01 00:06 UTC, reviewed episode-1-structure-v2-four-keys.md under the charter, and returned **REWORK**. This file records what the run log preserves; the log truncates the final review, so the text below is **partial** and is marked where it cuts. The full review is on the run session: claude.ai/code/session_01Q4zqn8hmkeYiW6UVSiVTSS. Stephen also received the push notification. If the full text is wanted on file, paste it and this doc gets it verbatim.

The three findings, and their disposition

#	Finding	Disposition
1	The structure's own attack round has not happened: only the attack prompt is committed, no result file exists, contradicting the run briefing's claim that the structure "already survived" adversarial review	Correct, and the briefing error was Claude's (the 2026-08-31 routine-prompt refresh overstated; corrected same day in the routine). The attack was already queued as the next step; finding changes nothing in the plan except urgency
2	Likely kill-gate failure on agency: all four decisions are Class A/B by the document's own words ("the ending-changing axis is unavailable"), against Stephen's standing ruling that decisions must change play	Escalated into the attack: the attack prompt gains front 5, a per-decision cosmetic-or-real test plus the design question of whether the unruléd close (Part F question F2) can carry a genuinely divergent decision without touching the settled spine. Final disposition is Stephen's, at the F2 ruling
3	Zero tester validation of the entire pivot: three readers said yes; the twelve have seen nothing, and the 12/12 history says internal confidence is not evidence	Stephen's call. Candidate cheapest test: the cold open v0.4 to the twelve (story goes to testers; mechanics stay self-verified per the standing tester ruling)

The push notification the run sent

Critical review of Four Keys structure v2.0 done: likely kill-gate failure on agency (all 4 decisions look cosmetic per the doc's own words) plus zero real tester validation of the whole pivot. Verdict: REWORK.

The review, as preserved in the run log (PARTIAL: the log truncates)

Critical Review: The Friends with Four Keys, Structure v2.0

Target: SAS/episode-1-structure-v2-four-keys.md (100 lead-packages / 10 chapters, committed 2026-08-31): the newest, largest, and most expensive-to-reverse deliverable in the last three days of work, and the one whose own attack round has **not actually happened yet** (only the attack *prompt* is committed; no four-keys-structure-attack-chatgpt-* result file exists). That contradicts this run's briefing that structure "already

survived" adversarial review: it hasn't. That gap, plus what three completed attack rounds (premise, spine x3) never checked, is where this review spent its time.

1. Objective and success criteria

Structure v2.0 exists to turn the ruled premise/spine/economy into 100 buildable lead-packages that (a) conform to the ruled economy (1,600 T1eq, ~6,880 CC), (b) deliver a real belief-turn every chapter, (c) clear the four kill-gates the project re-derived after the Listener's 12/12 rejection, and (d) give engineering enough to cost the still-unbuilt package system. I judged it against those four things, not against its own internal rubric.

2. Verdict up front

REWORK. Two kill-gates are, on the document's own evidence, either failed or untested by the one instrument the project's own post-mortem says is decisive. The arithmetic and self-audit inside the document are genuinely excellent: that is not the problem. The problem is what the excellent arithmetic is standing on.

3. What is working, and why

- **The economy conformance is real, not asserted.** Part E1 recomputes T1eq from Part B line-by-line and lands on 1,600 exactly, 100 packages exactly, matching the ruled model. That's a rare case of a design doc's own numbers actually adding up when checked.
- **Unusual self-honesty.** E5 names its own three weakest beats before an attacker would; line 251 states outright that "the ending-changing axis is unavailable by construction" instead of hiding it; E3 flags a rea

[THE RUN LOG TRUNCATES HERE, ~9,800 characters remain unrecorded. Full text: the session link in the header. Known from the truncated remainder's opening: section 3 continues with the lab-heavy family-mix self-flag; the verdict line, confidence, evidence-that-would-change-it and the cheapest-practical-test are in the unrecorded portion.]