

ATTACK PROMPT: THE FOUR KEYS

STRUCTURE v2.0

2026-08-31. For ChatGPT. Paste the block below with *episode-1-structure-v2-four-keys.md* attached (or pasted whole). The spine is settled and out of scope; this attack is aimed at the structure's own claims.

You are the adversarial reviewer for a merge-puzzle true-crime mobile game. Attached is the Episode One structure: about 100 "lead packages" in 10 chapters, each package a bundle of 1 to 5 merge-board cards whose fulfilment shows one story beat. You have previously attacked this episode's spine three times; the spine PASSED and is now settled canon. **Do not relitigate the spine:** its facts, days, character locks, sweep and rulings (Parts marked as sourced from spine v1.4, the B6 rows, the A5 clock) are out of scope. Attack the structure, on four fronts, in this order.

1. The beat census: is every chapter's turn real? The document marks 18 packages ★TURN and claims every chapter changes what the player believes at least once. For each of the 10 chapters, test the marked turn(s) against this standard: state the belief a player holds before the package, the thing the package's board work and beat deliver, and the belief after. A turn fails if the belief after is the belief before restated, if the change arrives in dialogue that the board work merely tickets (a toll booth), or if the change was already available to the player in an earlier chapter. Also audit the opposite error: any package marked (t) texture that actually turns belief, which would mean the pacing analysis is wrong. Name every failure with its package id. The document itself names its three weakest beats in Part E5; check whether its self-diagnosis is honest or whether worse ones are hiding.

2. The economy conformance. Recompute Part E1's arithmetic from the Part B tables (T1eq of tier $n = 2^{(n-1)}$ per item, quantity multiplies, chapter sums, the 1,600 total, the 6,880 CC at 4.3 per T1eq). Then attack the two soft spots the document flags: the 223-card count against the model's 250-card working figure (does the heavier-card preference break the "mass spread across many small cards, no walls" claim anywhere specific? name packages), and the Part E3 lab-heavy family-mix finding ($A+F = 63$ per cent of mass; is the stated $\approx 1,210$ -tap consequence computed right, and is "name it for the tuning pass" an adequate disposition or does it invalidate the 8-to-10-session guide now?).

3. The variety evidence. The ruling mandates variety: never a mechanical size progression, quick wins beside heavy builds late, a spike early, sizes and tiers mixed within every chapter. Part E2 offers evidence. Attack it: compute per-chapter package-size and T1eq sequences from Part B and test whether any chapter is in practice monotonic, whether the "quick wins" late are numerous enough to register (or token singletons), and whether chapter 2's ruled spike is actually a spike relative to its neighbours.

4. Chapter 1's pacing. The cold open v0.4 (~470 words) is spread across ten FTUE packages, first required tap proposed at ~0:12 to 0:15, chapter total 35 T1eq. Attack: does a ten-way cut fragment the open below narrative coherence (each package's VO stretch is listed in Part D; judge each boundary as a listening experience, not a word count)? Is a 12-second first interaction better or worse for a cold install than the 35-to-40-second band the project has been debating, given the open keeps rolling between merges? Is any of chapter 1's ten packages a beat a bored installer would quit on, and which?

5. The agency gate, promoted to a full front (added 2026-09-01 after an independent critical review flagged it as a likely kill-gate failure). The project's standing ruling: player agency must be integral, multiple decisions that change play, never a single end-of-episode gate; and the previous pilot died 12 of 12 with testers partly on this axis. The decisions (Part C) knowingly trade the ending-changing axis for who-is-protected/exposed/believed, because the spine's clock is settled. Attack each of D1 to D4 as a player would experience it: name what observably changes in play on each branch (different packages? different beats? different art, lines, flags with visible consequences later?), how many packages later the change is felt, and whether a player who chose the other way would ever know. A decision whose branches differ only in wording is cosmetic; say so per decision. Then answer the design question directly: within the settled spine, where could a decision with genuinely divergent play live, and is the still-unruled close (how it ends and what it costs, question F2) sufficient to carry one? Do not propose changing the spine.

Out of scope: the spine and all its facts; the v1.1 economy rulings (1,600 T1eq, ~100 packages, ~10 chapters, no optionals, CC 4.3 rate: these are Stephen-ruled); the redaction lines' content; prose quality (no prose exists yet).

Return: a verdict per front (PASS or REWORK), findings numbered and severity-ranked, each finding naming the package id or table row it attacks and the concrete repair it demands. No compliments. If a front passes, say so in one line and move on.

No em dashes in your reply.