

INDIGO CHIMP STUDIOS / CONTROLLED NARRATIVE PROMPT

ALLY QUINN

EPISODE ONE STORY

GENERATION AND

FALSIFICATION PROMPT


Research first. Falsify before developing.

VERSION 2.0

AQ-NAR-PROMPT-EP1-CHALLENGERS

24 August 2026 / READY FOR EXECUTION

MAJOR REVISION: OPERATIONAL REALITY IS NOW KILL GATE ZERO



Document control

Field	Controlled value
Document ID	AQ-NAR-PROMPT-EP1-CHALLENGERS
Title	Ally Quinn Episode One Story Generation and Falsification Prompt
Version	2.0
Date	24 August 2026
Status	READY FOR EXECUTION
Owner	Indigo Chimp Studios
Change class	Major revision
Supersedes	The earlier unversioned challenger-generation prompt used on 24 August 2026
Execution purpose	Generate up to three believable, murder-led Episode One challengers and attempt to disprove them before recommendation

Revision history

Version	Date	Status	Change summary
1.0	24 Aug 2026	Retired	Initial unversioned workflow. It tested story architecture but did not make operational plausibility a precondition. It promoted The Wrong Name despite a premise-dependent mortuary process being unverified.
2.0	24 Aug 2026	Current	Adds Kill Gate Zero, primary-source process validation, authority and custody mapping, competent-professional objections, antagonist least-cost testing, a no-quota rule, separate falsification, premise graveyard and explicit UNKNOWN-equals-FAIL governance.

Change-control rule

Major versions change gates, workflow or output authority. Minor versions add constraints or sources without changing the decision process. Patch versions correct wording only. Never overwrite a controlled version. Issue a new filename and revision row.

Known failure that triggered v2.0

The Wrong Name allowed a body-transfer contractor to remove a suspicious body and drive immediate cremation through forged order forms. In ordinary practice, police and coroner authority, custody, examination and release controls would intervene. The hook was developed before that mechanism was validated, and a foundational failure was carried as a later caveat instead of an automatic rejection.

Control principle

A compelling hook never earns permission to invent institutional powers. Documents record or communicate authority. They do not create powers the issuing actor does not possess.

Contents

Document control and revision history	2
How to use this prompt	4
Execution prompt: task and governance	5
Authority hierarchy and exclusions	5
Occupied and excluded story territory	6
Mandatory staged workflow	6
Kill Gate Zero: operational reality	7
Remaining kill gates	9
Agency gate and branch-cost control	10
Ally, cast, voice and prose controls	11
Creative lenses and season controls	11
Independent kill brief	12
Required execution output	13
Final execution safeguards	14
Execution record and future change control	15

How to use this prompt

Use the execution prompt in this document as a complete instruction set. Attach the six current project authorities named in the source hierarchy. Do not attach the retired pitch pack as creative inspiration. The three failed challengers are exclusion examples, not seeds to repair.

What v2.0 can and cannot promise

This process can greatly reduce believable-process failures and prevent an unverified premise from being recommended. No prompt can guarantee that every factual assumption is correct or that a story is superior before specialist review and player testing. The required safeguard is that uncertainty stops advancement instead of being hidden in a footnote.

Expected execution result

- A target of three new challengers, but fewer if fewer survive.
- A documented attrition path showing why near-miss premises were killed.
- Authoritative research for every premise-dependent institutional process.
- Three complete pitch packages only if three genuinely survive.
- An intrinsic forced ranking, followed by a comparison with G, H and I based only on supplied material.
- A recommendation of exactly one new table-read candidate or none.

Recommended operating sequence

- 1 Run the prompt without supplying your preferred incumbent ranking.
- 2 Inspect the reality dossiers before reading the elevator pitches.
- 3 Reject any finalist whose premise still relies on an UNKNOWN process.
- 4 Run surviving cold opens blind. Do not show titles, rankings or premise explanations.
- 5 Prototype the agency choices separately. Cold-open pull does not validate player agency.

BEGIN EXECUTION PROMPT

Core task

Develop **up to three** genuinely distinct, believable, murder-led Episode One challengers for Ally Quinn: True Crime Merge. The target is three. Do not manufacture three survivors. If reality, story or production gates leave fewer than three, output fewer and explain the attrition.

Your roles

Act as a senior crime-thriller architect, interactive narrative designer, procedural researcher and constructive adversary. These roles are sequential, not blended. Research and falsification must happen before creative development.

Success definition

A successful execution is not one that produces three polished pitches. It is one that refuses unbelievable premises early and produces only candidates whose core events, institutions, evidence access and antagonist behaviour survive scrutiny.

- Believable does not mean documentary-dry. It means the extraordinary events arise from credible human actions, authority, access and consequences.
- Superior means a candidate offers a stronger combined case for compulsion, human consequence, fair-play investigation, player agency, Havenbay necessity and production reality than the available incumbents.
- Cold-open compulsion is always UNKNOWN until tested.
- No scores or percentages. Use PASS, FAIL or UNKNOWN and forced rankings without ties.

No-quota rule

Never lower a gate, soften a failure or carry a premise-dependent UNKNOWN as a caveat merely to return three stories. UNKNOWN on the foundation of a story is FAIL for finalist selection.

Authority hierarchy and exclusions

Read these authorities completely and in this order

- 1 pivot-brief-2026-08-23.md. This is the current-state authority and wins every conflict about what is true now.
- 2 critical-reviewer-charter.md. This governs assessment behaviour and constructive adversity.
- 3 episode-1-vo-lock-v1.1-WIP.md. Only locked L1 to L5 dialogue is final voice evidence; its conventions govern game-facing prose.
- 4 ally-quinn-bible.md. Use subject to the pivot brief's explicit supersessions.
- 5 canon-character-roster.md. This controls character existence, status and names.
- 6 portrait-prompts-canon-cast.md. This controls approved, proposed and prohibited portrait facts.

Do not read or use

- The full Episode One pitch PDF. Incumbents may be compared later using only their supplied one-paragraph summaries and occupied-territory list.

- stage-backgrounds-ep1-kit, item-icon-regeneration-kit or art-kit-items-and-ui. They are stale art-pipeline documents and must not influence story creation.
- The retired challenger pitch pack as a source of ideas. The Wrong Name, The Morning Edition and The Last Carriage are failed or unvalidated examples, not repair briefs.

Source discipline

- Separate every claim into CANON, VERIFIED REAL-WORLD PROCESS, PROPOSED CANON or UNKNOWN.
- Do not invent economy figures, legal powers, police access, medical powers, forensic certainty, technical capabilities or production assets.
- Where an answer depends on material not provided, say exactly what is missing.
- Where a process has a realistic chance of being recalled incorrectly, research it before using it.

Occupied and excluded story territory

Existing occupied engines

- **The Listener:** nightly voicemails; wordless or interrupted recording; ambient audio locates a home; the woman is gone before the episode begins.
- **G - The Confession:** a confession recorded by a trusted mentor; the mentor becomes last witness and target; a 1989 harbour death was staged as a drunken fall; paid silence, payer voice and a keeper with decades of receipts.
- **H - The Overnight Girl:** childhood escape plan; coded phone ritual as dead-man switch; body in cold storage; a patient professional works through a trust network; breadcrumbs keyed to the show archive.
- **I - The Last Caller:** Tip Line confession cut off mid-name; room-tone forensics identify the caller; a childhood witness, now grown, is hunted.

Retired new-pitch territory

- **The Wrong Name:** a body carrying a living person's identity; contractor-controlled body transfer; administrative death through order forms. Do not repair, reskin or relocate this engine.
- **The Morning Edition:** murdered printer; advance obituaries for living owners; obituary and property-transfer machinery; Ally's premature obituary. Do not assume the mechanism can be fixed.
- **The Last Carriage:** murdered transit auditor; impossible passenger records; cloned taps; witness pursued through a closed line. Do not reuse the engine without a new, independently verified premise, and do not present a repaired version in this execution.

Structural exclusion

Do not use a damaged or interrupted recording reconstructed through audio analysis to locate an already missing person. This remains excluded for ideation diversity, not because it is proven to cause failure.

Mandatory staged workflow

Do not skip stages

Premise generation, process verification, falsification, story development and ranking are separate passes. Do not write cold opens, titles or emotional selling copy until a seed passes Kill Gate Zero.

Stage one: broad seed field

- 1 Generate at least twenty-four one-sentence murder-premise seeds internally.
- 2 For each seed, state the single real-world process on which the hook depends.
- 3 Reject occupied, retired, supernatural, child-harm, gore-led, coincidence-led and decoratively Havenbay ideas.
- 4 Do not name, title or emotionally embellish seeds yet. Attractive presentation must not influence viability.

Stage two: Kill Gate Zero research

Apply the full operational-reality gate to every plausible seed. Use authoritative, preferably primary, sources for premise-dependent legal, police, coroner, medical, forensic, transport, telecommunications, property, banking, employment, insurance, digital-security or regulatory processes.

If Havenbay's jurisdiction is not canonical, do not power a premise with a narrow jurisdiction-specific loophole. Use a broadly shared operational baseline, or identify the required Havenbay canon ruling and fail the seed pending approval.

Stage three: controlled expansion

- 1 Expand only seeds that pass Kill Gate Zero into victim, antagonist, motive, method, evidence and jeopardy summaries.
- 2 Apply the remaining story, agency, canon and production gates.
- 3 If fewer than three survive the first twenty-four seeds, generate one additional field of twelve seeds and repeat the gates once.
- 4 If fewer than three survive after that second field, stop. Output the smaller survivor set and the reasons. Do not attempt a third field in the same execution.

Stage four: independent falsification

For each apparent finalist, write a kill brief before writing its pitch. Review the mechanics without title, elevator pitch or emotional framing. Attempt to eliminate the story from four perspectives: competent professional, rational antagonist, ordinary true-crime player and interactive-narrative producer.

Stage five: finalist development

Only after a seed survives falsification may you write its blind cold open, one-page pitch and complete episode architecture.

KILL GATE ZERO: OPERATIONAL REALITY

Automatic-fail rule

A seed cannot advance while any premise-dependent authority, custody, access, procedural, technological or institutional claim remains UNKNOWN. Verify it, replace the mechanism or kill the seed.

Required first-five-event audit

Map the first five causal events, not merely five scenes. For each event provide:

Required field	Question that must be answered
Actor	Who performs the action?
Authority	What legal, professional, contractual or physical authority permits it?
Custody	Who controls the body, person, scene, evidence, record, vehicle, account or system at this moment?
Normal response	What would a competent professional ordinarily do next?
Safeguard	What rule, verification or colleague would obstruct the antagonist?
Bypass	How does the antagonist plausibly bypass that safeguard using access they actually possess?
Consequence	What risk or trace does the bypass create?
Evidence source	Which authoritative source supports the operational claim?

Any blank, circular or invented answer is FAIL.

Institutional-power rules

- Corruption can misuse an existing power. It cannot create a power the institution or actor never had.
- A contractor may perform a delegated physical task only after valid authority is granted. A contract does not automatically displace police, court, coroner, medical or regulatory control.
- A forged document may deceive a plausible recipient for a plausible period. It does not automatically neutralise independent checks, custody records, signatures, witnesses or system alerts.
- A database entry can mislead action. It cannot make a living person legally or physically inaccessible without a credible downstream mechanism.
- Police, coroners, doctors, forensic staff, journalists, archivists and digital specialists remain competent. Their competence must create constraints the antagonist has to overcome.
- Ally's PI licence does not grant police powers, unrestricted crime-scene access, medical records, protected digital systems or custody of evidence.

Commonsense challenge

Answer each question in one plain sentence. If the answer requires exposition, coincidence or 'because the plot needs it,' fail the seed.

- Why do the police not prevent this?
- Why does the coroner, doctor, court, employer, bank, platform or regulator permit this?
- Why is Ally present and allowed to do what the story requires?
- Why does the antagonist choose this method instead of a simpler and safer one?
- Why does the victim or keeper not take the obvious protective action?
- What would an informed viewer object to in the first story beat?

Antagonist least-cost test

- Why murder?
- Why now?
- Why this victim?
- Why this method and staging?
- Why not intimidation, theft, flight, bribery, destruction or silence through an easier route?
- What does the method expose, cost or risk for the antagonist?

If the antagonist is performing an elaborate thriller mechanism only because the player needs a puzzle, fail the seed.

Research standard

- Use official agencies, legislation, regulations, courts, professional bodies, technical standards, manufacturer documentation or peer-reviewed research where applicable.
- Use news, blogs, entertainment references and general summaries only to locate primary material, not as final support for a premise-dependent claim.
- Cite the exact source near the claim it supports. Do not provide a decorative bibliography detached from the logic.
- State where the source jurisdiction differs from Havenbay and whether the difference matters.
- Do not quote long passages. Summarise the operative fact in plain language.

Remaining kill gates

Gate one: immediate compulsion

- The episode is murder-led, restrained twelve-plus and recognisably true crime.
- The cold open contains a concrete contradiction, danger or irreversible event the player understands at listening speed.
- Cold-open compulsion is labelled UNKNOWN until blind testing.
- The first beat must survive the ordinary-viewer objection before it is polished for drama.

Gate two: two-track jeopardy

- There is present danger to a living victim, witness, keeper or intended target.
- There is material danger to Ally, not embarrassment, delay or inconvenience.
- Judge the case by the weaker track.
- State one central jeopardy sentence containing the closing window.

Gate three: medium-safe clock

- The clock is day-stamped, lead-state or genuinely session-contained.
- Regeneration can take one hundred and fifty seconds. Pausing and closing the app must not make the fiction absurd.
- No continuous live chase may secretly wait while the player merges for half an hour.

Gate four: action-changing midpoint

The midpoint must change what the player does next. A new interpretation without a changed action, route, target, risk or evidence strategy is FAIL.

Gate five: fair-play evidence

- Every conclusion is traceable to evidence actually obtained.
- The murder method, culprit and motive have separate proof chains where appropriate.
- A fair early plant exists before its meaning is explained.
- No coincidence proves the culprit. No withheld obvious fact manufactures a twist.

- The supporting cast remains competent and contributes domain-specific truth.

Gate six: human consequence

- The victim is grievable through one concrete human detail before becoming a puzzle component.
- A living keeper has a need beyond survival or delivering information.
- The criminal economy imposes a recognisable ordinary cost.
- The story does not use gore, child harm or melodrama to manufacture seriousness.

Gate seven: Havenbay necessity

Ask: could this exact case happen only in Havenbay? Atmosphere must advance evidence, access, motive, jeopardy or consequence. Fog, docks, old records and civic corruption are not sufficient by themselves.

Gate eight: production reality

- Fifteen to twenty played leads across two or three phases, determined by the story rather than a fixed table.
- The Field Kit is the starting generator. The Lab may unlock mid-episode. Kestrel Corner supports food, home and care functions.
- State played and total authored lead counts.
- Identify approved, proposed and new portraits and locations. Do not infer reuse from excluded art documents.
- Do not invent costs or economy figures.

Agency gate and branch-cost control

A single end-of-episode branch is an automatic FAIL. Each episode requires at least four decisions that change which leads spawn, what the player does next, what evidence survives, the risk state or the ending. Dialogue-only variation does not count.

Branch rules

- Every branch rejoins within a maximum of three leads.
- No nested fork opens before the current branch rejoins.
- At least two choices are Class A.
- No more than one choice is Class C.

Class	Authoring definition	Allowed effect
A	No exclusive voiced scene	Lead order, evidence quality, access, target protection, route or risk changes through shared assets
B	Two to six variant voiced lines	Shared scene with altered consequence, trust, evidence, risk or ending state
C	One fully exclusive lead or short chain	A materially different route that rejoins within three leads

Agency falsification questions

- Would a rational player consistently choose one option because it is obviously superior?
- Does the choice alter play, or merely choose the order of two mandatory scenes with no consequence?

- Can the player predict the type of trade-off without knowing the correct answer?
- Does at least one decision force a genuine editorial, evidentiary, protective or investigative sacrifice?
- Does Ally's irreversible ending change arise from player action rather than a final dialogue preference?

Ally, cast, voice and prose controls

Ally's required episode arc

- Player-facing podcaster first; licensed PI as a capability, not promotional identity.
- Early warmth that costs time, access or tactical advantage.
- Early competence visible through an action, not a speech about competence.
- Evidential discipline: distinguish fact, hypothesis and inference.
- A vulnerability, mistake or risk with a real consequence.
- An irreversible ending change caused by the player's investigation or publication choice.

Competent supporting roles

Character	Primary value
Del	Lawful action, institutional boundaries and defensible workarounds
Gerald	Pattern, police history and human judgement
Mo	Human knowledge and waterfront relationships
Arthur	Gazette archives and documentary memory
Vega	Digital investigation; pronouns remain deliberately unsettled unless a proposed ruling is made
Priya	Forensic trace analysis
Cortez	Cause and manner of death
Helen	Personal cost and family truth

Closed endearment canon

Gerald says 'love'. Mo says 'my lovely'. Dot says 'love'. Del calls Ally 'Quinny'. Nobody else gains an endearment.

Game-facing prose

- No em dashes.
- No exclamation marks.
- Grounded deadpan noir. Understatement over melodrama.
- Never explain the subtext.
- Every line parses at listening speed and stays near forty-five words or fewer.
- Numbers are spelled out in spoken lines.
- Use 'Grandad' with one d, and 'Okay'.
- Tip Line format is 'Message received: [weekday], [time]' with no calendar date.
- No monologues, exposition dumps or ornamental puns.

Creative lenses and season controls

Four development lenses

Lens	Question
Connelly	Can every conclusion be audited from obtained evidence, and do competent institutions create real constraints?
Harper	Could this case happen only in Havenbay, with place changing investigation rather than decorating it?
Hallett	What recordings, documents, images, manifests or objects does the player physically reconstruct?
Lehane	Who profits, who performs dangerous work, which institution absorbs the crime, who pays the ordinary cost and what compromises decent people?

Season-arc restraint

- The case is self-contained and resolved.
- Include one Thomas, Ferryman or Hart notch containing one new fact and one new cost.
- Do not unmask the Ferryman, depict a face, connect the episode killer casually to the season conspiracy or use a coin as a substitute for an actual link.
- The Episode Two handoff to The Ferryman must state the new fact, new cost, corrected misunderstanding and exact Episode Two use.
- Mark unsupported handoff material as PROPOSED CANON.

Independent kill brief

Before writing a positive pitch for a candidate, produce this adversarial assessment without title or selling language.

Professional objection

- Which professional would encounter the premise first?
- What would that person say is wrong, missing or impossible?
- Which normal safeguard does the story currently ignore?
- What evidence or authority would actually be required?

Rational-antagonist objection

- What simpler method achieves the same criminal goal?
- Why has the antagonist not already succeeded?
- Which step creates unnecessary exposure solely to give the player a clue?

Ordinary-player objection

- What is the first 'that would never happen' reaction?
- What is confusing before it is intriguing?
- Which question would make the player stop trusting the story rather than keep investigating?

Producer objection

- Is the hook hiding excessive location, portrait, VO or branch burden?
- Does the agency require authored material most players will never see?
- Does the evidence system remain readable on a mobile screen and at listening speed?

Decision rule

The kill brief ends with PASS or FAIL and the single decisive reason. No conditional PASS. A candidate may be revised once only if the repair preserves the core premise and replaces no verified fact with invention. Otherwise kill it.

Required execution output

Part one: source and attrition report

- 1 Confirm the six authorities were read in order and the excluded documents were not used.
- 2 State the number of seeds generated in each field.
- 3 Show the five closest rejected seeds and the decisive reason each failed. Do not beautify or rescue them.
- 4 List every real-world process researched for the survivors, with authoritative links beside the supported claim.
- 5 Provide a kill-gate table using PASS, FAIL and UNKNOWN. UNKNOWN on Kill Gate Zero is finalist FAIL.

Part two: one-page pitch for each survivor

Each begins with this exact heading:

ELEVATOR PITCH: WHY THIS SHOULD BE THE GAME'S EPISODE ONE STORYLINE

Then provide the high concept, opening promise, player activity, human engine, agency promise, Havenbay necessity and strongest honest reservation. Do not claim tested compulsion.

Part three: full package for each survivor

Section	Required content
A. Reality dossier	Verified baseline, first-five-event authority and custody map, antagonist bypass, sources, professional objection and remaining non-foundational unknowns
B. Identity	Title, high concept, dramatic question, player fantasy and exclusion statement
C. Blind cold open	Sixty to seventy-five seconds; no title, pitch letter, proper names or context briefing
D. Complete truth	Victim, killer, motive, method, staging, current activity, critical clue and fair early plant
E. Human engine	Grievable detail, keeper need, warmth cost and emotional reversal
F. Jeopardy and clock	Living-target track, Ally track, weaker-track verdict, closing window and pause-safe design
G. Evidence chain	Obtained evidence, supported conclusion, better question, wrong hypothesis, midpoint, proof and competent cast
H. Havenbay	Causal location and institution dependence; relocation test
I. Criminal economy	Profit, dangerous work, institution, ordinary cost and compromise
J. Phase topology	Two or three phases, fifteen to twenty played leads, total authored leads and generator progression
K. Agency	At least four decisions with class, consequence, divergence and rejoin
L. Ally, cast and art	Required Ally arc, cast purpose, portrait status and new or unverified production needs
M. Resolution and handoff	Self-contained resolution and precise proposed Episode Two bridge
N. Production reality	Relative burden without invented costs
O. Kill brief and verdict	Strongest professional, antagonist, player and producer objections; PASS or FAIL

Part four: ranking and slate decision

- 1 Rank the survivors intrinsically first, without incumbent comparison. No ties.

- 2 Force-rank every major dimension without aggregation.
- 3 Compare only the surviving candidates with G, H and I using the supplied summaries and occupied-territory list.
- 4 Recommend exactly one new table-read candidate or none.
- 5 State the strongest case against the recommendation and the evidence that should overturn it.

Final execution safeguards

Do not confuse polish with proof

Titles, cold opens, noir language, emotional details and evidence-board elegance are downstream. None may be used to argue that an unverified mechanism is plausible.

- Do not recommend a premise merely because its opening image is memorable.
- Do not label foundational research as 'to be validated later.' Validate first or reject.
- Do not repair the same failed mechanism with a different institution, contractor, document or city department.
- Do not infer that corruption makes normal controls disappear. Show the specific authorised person, access path, bypass and trace.
- Do not let Ally enter protected spaces, obtain private records or direct police action without a credible route.
- Do not make the competent professional conveniently absent, slow or incurious.
- Do not let a piece of evidence prove more than its real-world nature permits.
- Do not use a proposed Havenbay rule to rescue a seed during the same execution. Return it for owner approval first.
- Do not guarantee superiority. Recommend testing only after a candidate survives every non-testing gate.

Final pre-delivery checklist

Check	Required state
Authority and custody	Every first-five-event row answered and sourced
Professional reality	No unresolved informed-professional objection to the core mechanism
Antagonist rationality	No simpler obvious action defeats the premise
Ally access	Every protected action has a credible route
Cold open	Clear at listening speed; compulsion labelled UNKNOWN
Jeopardy	Both tracks material; lower track survives
Clock	Pause and regeneration safe
Midpoint	Player action changes
Agency	Four or more material decisions; branch limits obeyed
Evidence	Fair, traceable and no magical certainty
Havenbay	Case breaks or materially changes if relocated
Output quota	Fewer than three returned if fewer than three survive

Final instruction
Begin with source reading and seed falsification. Do not draft a cold open until the premise has passed Kill Gate Zero. If the final answer contains no surviving pitch, that is a valid and preferable result to an unbelievable Episode One.

END EXECUTION PROMPT

Execution record and future change control

Complete this record for each formal run. It separates the controlled prompt from the evidence produced by a particular model, source set and research session.

Execution field	Record at run time
Execution ID	Date plus sequential run identifier
Model and mode	Model name, reasoning mode and whether internet research was available
Project-source versions	Exact filenames and versions of the six authorities
Research record	Authoritative external sources used for premise-dependent processes
Seed field	Number generated in field one and, if invoked, field two
Gate Zero attrition	Number rejected for authority, custody, access, procedure, technology or antagonist-rationality failure
Final survivors	Zero to three candidate titles
Table-read recommendation	Exactly one candidate or none
Owner disposition	Accept, revise, reject or request a new controlled prompt version

Issue a new prompt version when

- A current-state authority changes a rule embedded in this prompt.
- A new failure exposes a missing or weak gate.
- An occupied-territory ruling changes after a pitch is approved or retired.
- Production constraints materially change lead count, branching, generators, cast or location limits.
- The output contract or recommendation authority changes.

Preserve the audit trail
Do not silently amend this PDF or reuse the v2.0 filename for changed instructions. Create the next controlled version, record the reason in the revision history and retain this file as the superseded baseline.