

PROMPT KIT v2.1 — PROPOSED CHANGE LIST

Against: [SAS/prompt-kit-ep1-challengers-v2.0.pdf](#) (AQ-NAR-PROMPT-EP1-CHALLENGERS v2.0) · **Date:** 2026-08-24 · **Status:** DIFF FOR REVIEW, not a controlled document

Governing source: [SAS/backbone-consensus-v1.3-FINAL.md](#), consensus reached, all clauses undisputed.

How to read this. Numbered, one item per change, each naming the v2.0 section it touches and the action taken. Nothing here is a rewrite. Where wording is short it is given in full; where it is long it is referenced to the consensus document by section, so the two cannot drift apart. **Twenty-eight items. Nine are structural, the rest are conforming.**

Review request: contest any item by number. Once the list is agreed, one party issues the controlled PDF and the other reviews it before it is used.

Part A — Purpose and output contract

These four are the reason v2.1 is a major version rather than a patch.

1. Core task — REPLACE

v2.0: "Develop up to three genuinely distinct, believable, murder-led Episode One challengers."

Reason: the job has changed. The immediate need is nine premise sketches, three per family, to feed the desk gate and the spoken-opening test. Generating full packages first would mean writing fourteen sections for premises that will be discarded at a five-question gate.

v2.1: two output modes, selected at invocation.

Mode A, premise sketches. Produce nine premise sketches, three per case family, each sufficient to answer the desk gate and no more. No cold opens, no titles, no elevator pitches, no full packages.

Mode B, full pitch package. Produce the complete package for a named premise that has already survived Mode A and the desk gate.

Mode A is a funnel stage, consistent with the staged workflow v2.0 already enforces. Mode B is v2.0's existing part three, conformed by the items below.

2. Mode A family requirement — ADD

Mode A must produce exactly three premises in each of: **fresh murder · never solved or misclassified · officially solved but contested.**

Reason, and this is the anti-bias instruction: the spoken-opening test is designed to falsify the backbone hypothesis. A prompt written under that hypothesis will produce weak fresh-murder premises by default, and a rigged slate produces a worthless result. **The fresh family must be generated with the same effort and, per item 3, first.**

3. Mode A generation order — ADD

Generate the fresh-murder family first, while fluency is highest. Generate the contested-record families last.

Reason: the mirror of the blind craft control in the consensus experiment. Authorial warm-up favours whatever is written last; the test's integrity requires the disfavoured family not to be written cold at the end.

4. No-quota rule — AMEND to apply per family

v2.0: return fewer than three overall rather than soften a gate.

v2.1: the rule now binds **per family**.

If a family cannot supply three premises that pass every gate, return fewer for that family and say so plainly. **If the contested-record families cannot each supply at least two, that is the desk failure the consensus document names, and it is reported as a result rather than repaired.**

Reason: the failure of a family to produce viable premises is the cheapest possible falsification of the backbone, and it happens before any tester is in the room. v2.0's global quota rule would have hidden it by topping up from a strong family.

Part B — The backbone, added as a tested hypothesis rather than as canon

5. Backbone frame — ADD, new section after Core task

The backbone under test. Internal engine: *the record is wrong, and somebody is still living with the consequence*. Player-facing promise: *Ally Quinn reopens the cases Havenbay has decided are over*. Episode construction rule: *a contested record, a live human consequence, rational resistance, reconstructive play, consequential agency, and an accountable resolution*.

This is not canon. It is a hypothesis awaiting the spoken-opening test. Generate under it, and generate the fresh-murder family as though it were false.

Reason: v2.0 has no concept of a series backbone at all. Adding it as canon would prejudge the test the kit exists to feed.

6. Six mandatory elements — ADD

Verbatim from consensus §2. Contested record · live human consequence · rational resistance · reconstructive play · consequential agency · accountable resolution. Any one missing is a REWORK.

Note: for a fresh-case premise, "contested record" may be an account that is *forming* rather than one already settled. Everything else applies unchanged. This is the one place where the fresh family is not held to identical wording, and it should be watched, because it is where a rigged comparison would enter.

7. Fresh cases — ADD

Fresh cases are permitted, not preferred, and carry **no numeric quota**. A fresh case satisfies the same six elements, the same access map, the same jeopardy gate and the same agency gate as any other. The season mix is an outcome of testing and ranking, never a rule set in advance.

Reason: v2.0 is silent, which would let a reader infer either a prohibition or a preference.

Part C — Gates

8. Kill Gate Zero — UNCHANGED, stated explicitly

Operational reality, the first-five-causal-event audit, institutional-power rules, the commonsense challenge, the antagonist least-cost test and the research standard all carry forward **without amendment**. v2.1 states this in terms, so that a reader of the new document does not assume the newer gates supersede it.

Reason: Gate Zero is the only gate with a proven kill record. It killed three premises in its first execution and two incumbents in its second.

9. Gate 2, two-track jeopardy — REPLACE

Replace v2.0's gate 2 in full with consensus §4 "Jeopardy gate". The material changes:

- **Human consequence** now requires a **named living person** and a **mechanism still running**, stated. A category fails. Historical suffering that has stopped fails.
- **Investigative consequence** replaces "material danger to Ally" with an enumerated list: evidence lost or made inadmissible · a witness permanently closed to the show · a hearing or prosecution prejudiced · a person exposed before they were ready · a remedy delayed past its opportunity · Ally's ability to be trusted, publish or work lawfully permanently reduced. Reversible-within-the-episode fails.
- **Legibility** is scaled to every episode, not only the pilot: cold open for the pilot, first podcast milestone otherwise.
- **Physical danger is one valid form of either track, neither required nor default.**

Reason: v2.0's "material danger to Ally" is the clause that generated contrived threats and, in the incumbent audit, two antagonist-rationality failures invented largely to satisfy it.

10. Gate 3, clock — AMEND

Retain the three structures. Add:

The **session-contained window is PARKED and not available for Episode One** per the ruling of 2026-08-24: a player who exhausts energy inside the window is either shown a broken fiction or gated out of a climax.

A clock is **optional**. It is not required where the current harm is already accumulating. If used it must be case-born: a scheduled hearing, transfer, lawful disposal, demolition, appointment, payment or official decision.

A deadline is never resistance. Clock and resistance are separate pitch fields and may not be satisfied by the same object.

11. Gate 4, midpoint — UNCHANGED.

12. Gate 5, fair-play evidence — AMEND

Add the access map's fifth field as a gate condition:

The pitch must state **what Ally gets wrong because she can only reach half the record**. The inaccessible half is where the wrong hypothesis lives.

Reason: v2.0 requires a plausible wrong hypothesis but supplies no method for generating one. Every pitch audited under v2.0 was marked down for lacking one, including all three of its own outputs. This makes the access map produce it.

13. Gate 6, human consequence — MERGE into gate 2 and DELETE as a separate gate

Reason: it now duplicates the human-consequence track. Its unique clauses (grievable detail before puzzle component; no gore, child harm or melodrama; keeper with a need beyond survival) move into gate 2 rather than being lost.

14. Gate 7, Havenbay necessity — UNCHANGED.

15. Gate 8, production reality — AMEND

Add to the required reporting: **reconstruction mode** (axis two) and **resolution mode** (axis three), so the season anti-formula audit has data at pitch time rather than after the fact.

16. Rational resistance — ADD as a named gate

Currently implicit in the antagonist least-cost test, which asks why the antagonist acts. Resistance asks a different question: what pushes back when the record is challenged.

Every episode requires an **active resistance system**: an agent or interest that benefits from, depends on, or is protected by the accepted record, and that responds to being investigated. A person protecting liberty, access, money, status, family or reputation · an institution protecting legitimacy, liability or authority · a witness or keeper protecting another person or a dangerous secret · a community enforcing a shared account · an evidentiary constraint created by lawful custody, degradation or prior error.

It may not be a deadline, a schedule or the passage of time. Its restraint must be as rational as its action.

Part D — Agency

17. Agency gate — REPLACE

Replace v2.0's "at least four decisions" section with consensus §4 "Consequence-and-agency gate" in full, including the material-loss clause bound to the end-state difference, the perceptibility clause, and the counterfactual closing test.

18. Fairness test — ADD

Consensus §4 "Fairness test", five statements required per consequential decision.

19. Branch-cost control — UNCHANGED, carried forward

At least four decisions · at least two class A · at most one class C · rejoin within three leads · no fork before rejoin · report played and authored lead counts separately. Consensus does not supersede these and they remain in force.

20. Reference case — ADD

Cite as the reference pattern: the player can protect a keeper or preserve decisive evidence, and cannot do both.

Reason: the four abuse patterns were adjudicated identically by both reviewers, and this one passed every clause without strain. A worked example prevents the gate being satisfied by paraphrase.

Part E — Taxonomy and anti-formula

21. Three-axis taxonomy — ADD, replacing v2.0's implicit single axis

Consensus §5 in full. Axis one, what is wrong with the record. Axis two, what the player reconstructs. Axis three, how the episode resolves. Plus the live-consequence list, which pairs with axis one and is **not** itself an anti-formula axis.

A pitch names one value on each axis. **Neither axis two nor axis three may repeat in consecutive main episodes.**

Reason: two episodes in different axis-one cells can still be sixteen leads of paper followed by an arrest. Axis one varies what the writer chooses; axes two and three vary what the player does and how it ends.

22. Wrongful conviction cap — ADD

Provisionally limited to one main episode per season: formula risk, research burden, and because it pulls Ally toward a lawyer's investigator when podcaster leads the player-facing identity.

23. Reconstruction requirement — ADD, sharpened

The merge requirements **are** the reconstruction, not the price of reaching it. Name the physical activity at pitch stage. A pitch that says the requirements "are" the investigation without saying what the player's hands do has stated an intention, not a mechanism, and fails.

Reason: verbatim the defect found in an incumbent, where "the merge requirements ARE the decoding" survived a full audit round before anyone noticed it described nothing.

Part F — Territory

24. Occupied territory — RESTRUCTURE

v2.0 lists The Listener, G, H and I together as occupied, and Wrong Name, Morning Edition and Last Carriage as retired. Three things have changed since.

Live candidates, do not duplicate: G "The Confession" (conditional pass, two repairs outstanding) · Slack Water · What the Log Says · No Fixed Address.

Failed at Kill Gate Zero, do not repair, do not reskin: H "The Overnight Girl" (antagonist search method; unexplained murder; disposal chosen for the audience) · I "The Last Caller" (a rational interrupter stops the call at sentence one; magical forensics on the caller's identity).

Retired, do not repair: The Wrong Name · The Morning Edition · The Last Carriage.

Demoted, not available as a pilot: The Listener, which holds no season slot until it re-passes the current gates.

Each entry keeps its device list from v2.0 so the anti-duplication boundary is unchanged in substance.

25. Structural exclusion — UNCHANGED

The interrupted-recording-to-locate-a-missing-person shape stays excluded, and stays labelled a **scoped ideation boundary rather than a proven diagnosis** of The Listener's failure. The wording in v2.0 is already correct and must not be hardened in v2.1.

Part G — Output and record

26. Mode A output contract — ADD

Per premise: one-paragraph premise · axis one, two and three values · the named living person and the mechanism still running · the resistance system · what the player's hands reconstruct · the first consequential decision and what it changes · the access map's five fields for the decisive item only · Kill Gate Zero verdict with the first-five-event audit · the five desk-gate answers.

No cold open. No title. No elevator pitch. No ranking.

Reason: ranking at sketch stage would prejudge the test.

27. Mode B output contract — AMEND v2.0 part three

Keep sections A to O. Add: the full access map for every decisive item, not only the first · axis values · resolution mode · the fairness test per decision · played and authored lead counts, which v2.0 already requires and which no output has yet supplied correctly.

28. Execution record and revision history — AMEND

Add fields: mode invoked · premises generated per family · desk-gate pass rate per family · axis values used · whether the fresh family was generated first. Add the v2.1 revision row recording that the change was driven by the backbone consensus rather than by a failure.

Items deliberately NOT changed

Recorded so their absence is not read as an oversight.

Kill Gate Zero in every particular · the staged workflow and its prohibition on writing a cold open before a premise passes · the independent kill brief written before any positive pitch, with its four objection frames and its no-conditional-pass rule · the research standard and its preference for primary sources · the source-discipline labels CANON, VERIFIED REAL-WORLD PROCESS, PROPOSED CANON and UNKNOWN · the season-arc restraint, one notch and one cost · the closed endearment map · every house prose rule · the four craft lenses · the final safeguards and the pre-delivery checklist.

Open question for the reviewer

One item I could not resolve alone. **Item 6's carve-out** lets a fresh-case premise satisfy "contested record" with an account that is *forming* rather than settled. That is necessary or the fresh family cannot exist at all under the six elements, and it is also the single place in v2.1 where the two families are not held to identical wording. If the test then favours contested records, someone will reasonably ask whether the gates were tilted.

I see two options and no clean third. Either accept the carve-out and record it as a known asymmetry to be reported alongside the test result, or drop "contested record" from the six elements for fresh cases entirely and require the other five, which is a larger asymmetry but an honest one. **I lean to the first and do not hold the view strongly.**