

PROMPT KIT v2.2 — PROPOSED CHANGE LIST

Against: [SAS/prompt-kit-v2.1-FINAL.pdf](#) (AQ-NAR-PROMPT-EP1-CHALLENGERS v2.1, READY FOR EXECUTION) · **Date:** 2026-08-24 · **Status:** DIFF FOR REVIEW, not a controlled document

Trigger: not a conformance failure. v2.1 was correctly implemented and correctly audited. It was then executed once, and the execution exposed a defect in the agreed design that four rounds of conformance auditing could not have found.

Twelve items. Three structural, the rest conforming. Contest by number.

The finding that triggers this version

Execution EX-2026-08-24-02 produced nine surviving premises across three families, with distinct record shapes, reconstruction modes and resolution modes. It satisfied every gate and every taxonomy rule in v2.1.

All nine share one moral mechanism. In every one, the institution behaves correctly and still produces the wrong answer. A fire service reaches an honest undetermined. A doctor lawfully certifies a death he did not examine. Samples are destroyed exactly on schedule. Expert testimony is correct by the standards of its year. Nine times, nobody did anything a listener could be angry about.

This was caused by the gates, not by the author

Kill Gate Zero's institutional-power rules forbid five things: corruption creating a power, a contract displacing control, a forged document neutralising checks, a professional being absent or incurious, a database making a person inaccessible. The antagonist least-cost test then eliminates anyone doing something elaborate.

Stack those and one configuration survives: **a small cheap act by one person, concealed by a system behaving normally.** The gates do not merely permit that shape. They funnel into it.

The general form of the defect: every gate in v2.1 is a prohibition, and a generator constrained only by prohibitions converges on whatever survives them. Nothing in the document requires variety of any kind that the three axes do not already name, and the three axes name categories of error rather than causes of it.

Why it is urgent rather than a season-variety concern

The nine feed the six-opening test. If all six openings share one moral shape, the experiment compares fresh-with-no-culprit against contested-with-no-culprit. Testers rejecting all six may be rejecting **the absence of anyone to be angry at**, and the result would be read as "no family pulls."

The panel already told this project once that a story was too safe. Nine premises in which nobody is culpable is the worst possible material to put in front of them, and a null result would be indistinguishable from a real one.

Nothing is lost if this lands before the openings are written, which is the next step.

Part A — Structural

1. Axis four: failure mechanism — ADD

v2.1 has three axes: what is wrong with the record, what the player reconstructs, how the episode resolves.

Reason: axis one names a **category of error**. It does not name the **cause**. Two premises can both be "misclassified" and be morally opposite depending on whether a doctor was inattentive or a brother lied. The observed monoculture sat entirely inside axis one's blind spot, which is why the taxonomy did not catch it.

Axis four, why the record is wrong. Name one at premise stage.

1. **Correct system, wrong answer.** Every actor performs correctly and the outcome is still false. Nobody is culpable.
2. **Misuse of a power actually held.** Not an invented power: a discretion, an access or an authority the person genuinely has, exercised for a reason. Explicitly permitted by Kill Gate Zero.
3. **Overruled by someone with standing.** The system worked and lawful authority set it aside.
4. **Unequal looking.** Nobody looked hard because of who the victim was. Not a failed process, an unequal one.
5. **A plausible person lied to the system,** and the system believed them, because that is what systems do with plausible people. The institution is a tool, not a failure.
6. **A handoff gap.** Two institutions each did their job and the space between them swallowed it.
7. **A community enforced the account.** Everyone knew, nobody said.
8. **The victim's own choice created the record.** They lied, to protect someone or themselves, and the account followed.

The list is not exhaustive. A mechanism outside it must be named, described in one sentence, and flagged as an addition in the execution record. See the open question below.

Note what several of these restore: a person the audience can be angry at, without requiring a conspiracy or an invented power. That is what the monoculture removed.

2. Field-level mechanism distribution — ADD

Reason: this is the item that actually prevents recurrence, and it is the first **positive requirement** in the document. A taxonomy that merely records a value does not stop nine premises sharing it.

Every seed field must exhibit at least three distinct failure mechanisms, and no more than one third of the surviving premises in a field may share one mechanism.

Applied to a nine-survivor field: at least three mechanisms present, and no mechanism on more than three premises. If the rule cannot be met without weakening a gate, **return fewer survivors and report the mechanism distribution as the reason.** The no-quota rule takes precedence.

3. Rotation extends to mechanism — AMEND

v2.1 rotates the starting family across three rounds. It says nothing about mechanism.

Reason: family rotation controls order advantage between families. It does nothing about the generator settling into one moral groove within them, which is exactly what happened.

Within each round, **no two consecutive seeds may share a failure mechanism**, across families as well as within them. Record the mechanism alongside each seed at generation, before any gate is applied.

Part B — Conforming

4. Consecutive-repeat rule extends to axis four — AMEND

Neither axis two, axis three **nor axis four** may repeat in consecutive main episodes.

Applies at slate construction, not at premise generation, consistent with the existing rule.

5. Mode A output contract — AMEND

Add per premise: **axis four value**, with one sentence naming who is culpable and whether the audience can reach them.

Add to Part One: a **mechanism distribution table** showing the count per mechanism across survivors, and an explicit statement of whether the item 2 rule bound.

6. Mode B output contract — AMEND

Section B, Identity, gains the axis four value. Section O, kill brief, gains the item 9 question below.

7. Execution record — AMEND

Add: **failure mechanism per premise · mechanism distribution across survivors · whether the distribution rule bound and what it excluded · any mechanism proposed outside the list of eight.**

8. What this process can and cannot promise — AMEND

Reason: the general lesson belongs where a future reader will meet it, and v2.1 already has a section admitting the process's limits.

Every gate in this document is a prohibition. A generator constrained only by prohibitions converges on whatever survives them. v2.1 was correctly implemented, correctly audited over four rounds, and produced nine premises sharing one moral mechanism, because the gates that secured believability silently selected for it. Conformance auditing cannot find this class of defect: it verifies that a document says what it was agreed to say, not whether the agreed thing has a blind spot. **Only execution finds it.** Where this document requires variety, the requirement is positive and stated, and future versions should add positive requirements rather than further prohibitions wherever a prohibition would do.

9. Kill brief, ordinary-player objection — AMEND

Reason: item 2 catches monoculture at field level. This catches it at premise level, costs one line, and would have flagged all nine individually.

Add to the ordinary-player objection: **Is there anyone in this story the audience can be angry at? If not, is that deliberate, and does the episode survive it?**

10. Gate 9 and axis four are distinct — ADD, one clarifying line

Reason: Gate 9 lists "a community enforcing a shared account" as a resistance form, and it is also mechanism 7. They are related and not the same, and merging them would let one field satisfy both.

Resistance is what pushes back when the record is challenged. Failure mechanism is why the record was wrong in the first place. A premise may share a value across the two, but it must state both separately and may not satisfy one by citing the other.

11. Consequence for execution EX-2026-08-24-02 — ADD

The nine surviving premises of EX-2026-08-24-02 predate axis four and all nine carry mechanism 1. They must be **mechanism-audited before any opening is written**. Premises beyond the item 2 cap are returned to Mode A for regeneration against the new axis. Those retained keep their Kill Gate Zero verdicts, which are unaffected.

Recommendation, for the reviewer to contest: retain three, regenerate six. Binning all nine discards verified Gate Zero work and research for no gain; retaining more than three breaches the rule this version exists to create.

12. Document control and revision history — AMEND

Version 2.2 · status READY FOR EXECUTION only after review · supersedes v2.1 · change class major · governing source **backbone-consensus-v1.3-FINAL** · revision row recording that this version was triggered by an execution finding rather than a conformance failure, and that v2.1 was correctly implemented.

Deliberately NOT changed

Gate 2, jeopardy. It requires a named living person under an operating consequence. It does **not** require a named culpable person, and it should not. Adding that would force culpability into every premise and abolish mechanism 1 entirely, which is a legitimate and often excellent shape. The monoculture is a problem of nine, not of one. The correct instrument is a field-level cap, not a gate.

Kill Gate Zero, in every particular. The institutional-power rules are why the premises are believable. They are also the proximate cause of the monoculture, and the fix is not to weaken them but to require variety alongside them. Nothing in Gate Zero is touched.

The antagonist least-cost test. Same reasoning.

Everything else in v2.1, including the staged workflow, the eight-per-family rule, the consequence-and-agency gate, the access map, the territory section and the prose rules.

Open question for the reviewer

Should the list of eight mechanisms be closed or open?

Item 1 makes it open, requiring any addition to be named and flagged. That preserves honesty about a taxonomy nobody has validated.

The argument the other way is stronger than I first thought. **An open list lets a generator invent a ninth mechanism to escape the item 2 cap.** Faced with "no more than three may share mechanism 1," the cheapest evasion is to relabel a fourth premise as a new mechanism that is mechanism 1 wearing a hat. A closed list is auditable; an open one is not.

Two options, no clean third. **Closed**, with additions permitted only by an agreed version change, which is rigid but enforceable. Or **open with a guard**: additions permitted, but a premise using a novel mechanism cannot count toward satisfying the distribution rule in the same execution that proposes it.

I lean to the second and hold it weakly. It is the reviewer's call, and it is the only item here I would not implement without one.