

OUR RULES UNDER EVIDENCE

What the genre research and the alignment audit actually change about AQ30, and what to do this week · 2026-08-27

Scope ruling recorded, Stephen 2026-08-27: FOUR EPISODES MUST SHIP FOR PRODUCTION. THE DATE CAN SLIP. Episode count is now the fixed constraint and 1 October is not. Everything below is written against that.

1. The verdict, in one paragraph

The audit says our framework is right and unproven, and it is correct on both halves. Kill Gate Zero, the access map, the agency gate, the clock convention and the evidential-language discipline all map cleanly onto what the research says actually drives the genre. None of them has ever been tested in play. The only direct player evidence this project owns is twelve people rejecting the one thing we shipped. **The risk is not that we built the wrong rules. It is that we built an exceptionally well-governed non-event,** and no amount of further auditing between models can tell us which.

2. What the evidence confirms, so stop re-litigating it

Kill Gate Zero and the least-cost test · professionals stay competent · the five-field access map including the inaccessible half that produces Ally's wrong hypothesis · rational resistance as distinct from a deadline · day-stamped and lead-state clocks and the ban on timers that punish absence · the consequence-and-agency gate with material loss and visible attribution · separating fact, claim, inference, allegation and unknown · no gore-led premises · the twelve-plus boundary · warmth and contrast through Gerald, Mo, Del and the Rusty Anchor.

Two of these are directly supported by the strongest evidence in the report. **Contrast is a real engine**, not a softening: dark material lands harder against warmth, which retires the idea that unrelenting dread would have fixed The Listener. And **participation is asymmetric** — one accurate tip can change a case while a thousand speculative posts harass an innocent person — which is exactly why our observation/inference/allegation discipline earns its cost.

3. Six of our own rules that the evidence challenges

Each is quoted from our documents, with the specific change.

3.1 The tone slogan is doing damage • [ally-quinn-bible.md:54](#)

Investigative, not brutal. "The darkness is intellectual intrigue, not bloodstains."

The research is blunt: **the meaningful distinction is not graphic versus non-graphic. It is whether a detail performs evidentiary, causal, emotional, historical or testimonial work.** Detail that changes understanding can be unflinching; detail that decorates cruelty turns a person into a texture.

Our slogan sets a *graphicness* threshold where the evidence sets a *function* test. It reads as permission to abstract bodily consequence away, and twelve testers told us the result felt too safe.

Change: keep the rating boundary, replace the test. A physical detail is permitted when it establishes mechanism, testimony or consequence, and refused when it exists for effect. That is a stricter rule than the slogan, not a looser one, and it stops us writing deaths nobody can feel.

3.2 The audience section states contested claims as settled • [ally-quinn-bible.md §1.4](#)

true-crime podcast audiences run ~73% female and highly engaged ... they "crave a sense of participation"

The female skew is real and well evidenced. **The single-motive explanation is not.** The research explicitly warns that the safety-tips account is incomplete, can place responsibility on women, and overstates stranger danger while obscuring violence in intimate settings. Motives are plural: threat learning, empathy, anger at institutional failure, social belonging, escape, entertainment, and the moral permission to enjoy dark material.

Change: restate §1.4 as plural motives with the market figures kept and their denominators noted. Pew's 34 per cent and Edison's ranking are not interchangeable measures and our doc currently treats them as one.

3.3 Marketing proposes exactly what the audit forbids • [ally-quinn-bible.md:926](#)

real *Echoes of Havenbay* teaser episodes on Spotify/Apple **blur game and reality**

Fiction cannot inherit the reality premium. True crime's voltage comes from the referential claim that this happened to a real person, and a fictional game cannot borrow it. Attempting to would spend trust we need for the Tip Line's own promise.

Change: cut the blur. Teaser episodes are fine and good; presenting them as actuality is not. The crime-doc trailer aesthetic survives untouched.

3.4 We have been writing mechanisms where the evidence wants minds

The strongest morbid-curiosity predictor of true-crime consumption is **interest in the minds of dangerous people**. Interest in bodily harm is much weaker and does not uniquely predict consumption once overlapping motives are controlled.

Look at what we have actually produced. Forty-five premises whose antagonists are overwhelmingly *situations*: a retention schedule, a misfiled record, a system behaving correctly. The one opening that drew the strongest craft response put a man on the phone to his girlfriend's employer sounding like he was cancelling a dentist appointment.

Change: no new rule. A generation instruction. The antagonist must be a person the audience gets access to, not a mechanism the player discovers. This is the same finding the participating-subject engine encodes and it is why Dean-being-helpful works.

3.5 The reason-to-care rule stops one level too high

The research's deepest line: *"It is most compelling when every new fact changes not only what the audience knows but what the audience believes the story is about."*

We require a one-sentence contested record and a one-sentence jeopardy at premise level. We require nothing at phase or lead level, which is where a player actually is.

Change, and it is editorial and cheap: one non-canon review field at game, season, episode and phase level — who may lose what, what is not known, why the next evidence can change it. If a phase cannot state its live question, its leads get combined, cut or reframed. Not a gate, a test instrument.

3.6 Our merge loop has never been shown to change anyone's mind

This is the audit's rank-one gap and the most serious item in either document. The bible says merges reconstruct evidence. The shipped loop requires coffee cups, bar items and tools, and records activated leads. **No test anywhere in this project establishes that the physical act of merging changes what a player believes.**

If it does not, then merging is a toll paid to hear the next scene, and every other rule we have is decoration on a payment mechanic.

4. Where I disagree with the audit's sequencing

The audit runs: blind opening test → greybox the winner → first-session test. Sensible for one pilot.

Do the evidence-turn test in parallel, not after, and do it now.

Its rank-one gap is a question about the *game*, not about the episode. Whether a merge can change a belief does not depend on which story we choose, and the audit's own falsification test for it needs no pilot at all: a paper or clickable prototype of one evidence turn, two competing explanations, ask which the player holds before and after.

That costs an afternoon. If it fails, the story selection matters far less than a loop redesign, and we would want to know that before writing four episodes. Running it after the opening test means discovering the loop problem having already spent the openings.

Recommended order: the six-opening screen and the one-turn evidence paper test run this week, independently. Then greybox the winner through one reconstruction and one decision. Then the first-session test.

5. What four episodes changes about all of this

It makes the testing more valuable, not less. A mechanism proved once is reused four times; a mechanism assumed is wasted four times. The case for greyboxing before content production roughly quadruples.

It moves the representation check from later to now. The research documents that the missing-white-woman pattern persists in leading podcasts, and that a work can call itself victim-centred while still selecting victims through an inherited hierarchy. We are about to choose four cases at once. Our current candidate is a woman in her thirties killed by a former partner; the retained seed is an unnamed woman under a viaduct. **If all four victims are young women and the danger is always an intimate partner, we will have reproduced the exact pattern the evidence names, and we will have done it deliberately because we selected the slate in one sitting.** Check the four together, before committing, not after.

It gives the three-axis taxonomy real work. Record shape, what the player reconstructs and how it resolves must not repeat consecutively. With one episode that rule was theoretical. With four it is the difference between a season and the same episode four times.

It changes nothing about the tests. Four episodes do not need four first-session tests. They need one mechanism proven and then applied.

6. What I would do this week

1. **Finish the six-opening screen.** Already running, costs nothing more.
 2. **Run the one-turn evidence paper test.** Two competing explanations, one reconstruction, before-and-after. An afternoon.
 3. **Make the four editorial changes** in section 3.1, 3.2, 3.3 and 3.5. These are stale or actively misleading documentation, not adoption of untested recommendations, and leaving them costs us whenever someone reads them as current.
 4. **Write the forty-eight hours** of the chosen premise, which was already the agreed next step and which every opening failed for want of.
 5. **Do not** freeze the hook field, write the twenty-four cards, spend the completion pass, commission VO, art, or economy regeneration.
-

7. The honest risk

The audit's strongest argument against itself is that adding test stages slows a project that is already behind, and that is real. But the count is now fixed and the date is not, which removes most of the force of that objection.

The bigger risk is subtler. **This project's failure mode is producing beautiful instruments instead of doing the work**, and a rules-audit of our rules-audit is exactly how that recurs. The two documents in front of us are good, and their single most useful sentence is the one that says: *a gate requiring a consequential question is not proof that a listener urgently wants its answer*.

Nothing in section 3 above is proof either. It is tidier documentation. **The only items on this page that can actually tell us anything are numbers two and four in section 6.**