

STORY PRODUCTION PIPELINE

How to produce quality stories with AI, efficiently, with no writing budget · 2026-08-27 · Derived entirely from what this project's own record shows, not from theory

The principle

Optimise for the shortest path to a human reaction, not for the fewest human interruptions.

Every quality gain this project has made came from Stephen or Trish reacting to something. None came from a gate. The loop we have been running takes about a week to reach one human reaction. The loop below reaches one in about a day, and the pre-flight exists only to make sure that reaction is worth their ten minutes.

Three findings from our own data, which override any framework:

1. **Depth beats breadth.** Forty-five premises generated wide produced one survivor. One premise developed deep produced the leading candidate. Stop generating fields.
 2. **The failures are a short recurring list,** not unpredictable quality. Thirteen objections cover essentially every rejection in three weeks, and the thirteenth was added by the rejection it caught.
 3. **Constrain the check, not the generation.** Prohibitions applied before writing produced conformant and contrived work. The same rules applied to a finished draft catch the same defects without deadening the prose.
-

The pipeline

One pass is about a day. Failures are front-loaded into steps where they cost nothing.

#	Step	Who	Output
1	The spine. What actually happened, hour by hour, and what Ally can lawfully reach, from whom, on which day	Claude	Timeline + access ledger. No prose.
1b	★ The earliest-action sweep. One page. Added 2026-08-27 and it is the most important step on this table	Claude	A first-action date per person, or a kill
2	Pre-flight. The thirteen objections applied to the spine, each answered in writing	Claude	Thirteen written answers or a kill
3	Adversarial audit. Attack the spine, not the prose	ChatGPT	Findings; repair or kill
4	Prose. Opening and key beats, written from the repaired spine	Fable	Draft copy
5	Reader check. Two questions, ten minutes	Stephen, Trish	Grab / believe
6	Kill or continue. Salvage fragments either way	Stephen	Decision

Step 1 is the one we have never done and it is why every opening failed. Six openings, three writers, and all six were holed on "how did she get that." Ally could not source what she knew because nobody had decided what she actually got. A spine written before any prose makes that class of failure impossible rather than catchable.

★ STEP 1b: THE EARLIEST-ACTION SWEEP

Added 2026-08-27, authored by ChatGPT after it found the same structural defect in three consecutive premises. It is not an objection and it does not belong in the pre-flight. It is a mechanical sweep, it takes one page, and it would have killed all three premises before an audit ever saw them.

The defect it catches

Accurate nodes and a false graph.

The engine gets declared first. Then every row of the access ledger is checked *locally* for lawful possession, and every row passes. **Nobody ever recomputes the earliest action available from the rows combined.** So the spine ends up full of people who can each lawfully act, silently prevented from composing into an intervention that would end the episode at lead 3.

Three premises died or reworked on this, each time inside the section written to prevent it:

Premise	The node that could act	What silently stopped it
No Fixed Address	A workplace supervisor, reachable and willing	The engine needed the identity route to be the only one
The Same House Twice v1.0	The certifying doctor, reachable on day three	He was simply not visited until lead 10
The Same House Twice v1.1	The same doctor, now visited on day three	His admission was declared worthless, and his live clinical duty to his own symptomatic patient was never mentioned

How to run it

Ignore the authored lead order entirely. Start at Day 1. For every person in the access ledger, ask:

Given everything this person can lawfully know today, what is the earliest competent action they can take to stop the harm or to trigger an authority?

Write down the first day of each action, in a table, with the person and the action.

If any first-action date precedes the planned turn, the spine fails. Not "has a note against it" — fails, before it goes to audit.

★ THE RECIPIENT COLUMN, added 2026-08-27 after the sweep was run and gamed

The first spine written under this sweep passed it, and the sweep was wrong. Not because the instrument is weak, but because it was run one column short.

The habit does not disappear when you stop silencing people. It moves into latency.

The initiators all acted correctly and on time. **The recipients were left outside the graph:** the care provider who receives the escalation, the authority that receives the doctor's statement, the office that receives the grave application, the pharmacy that holds the note. **"It takes weeks" described the outcome the story wanted, not the recipient's earliest action.** That is the same suppression re-expressed as delay, and it is harder to see because nothing in the document is silent.

So every row gets four more cells:

Action	Recipient	Recipient's earliest response	Every duty the combined facts trigger	Earliest consequential outcome
--------	-----------	-------------------------------	---------------------------------------	--------------------------------

Two rules for the new cells:

- Name the recipient. "The system" is not a recipient.** If nobody can be named, the latency is invented.

- **★ A slow process does not absorb a fast duty.** An application that takes weeks does not establish that the same information cannot trigger an immediate protective response about the person currently at risk. **Those are two different duties on the same facts and both must be run.**

Three rules that make it work

1. **Include professional duties, not just powers.** The question is not only what a person *may* do, it is what their job *requires* them to do once they know. A doctor told his own patient may be being overdosed does not have a choice about ringing her.
2. **A duty renews when the facts change.** A concern reported seven weeks ago does not discharge the duty created by a new deterioration this week.
3. **Every suppression must be named and paid for.** If a node genuinely cannot act, the reason goes in the sweep in one sentence, and it must be a reason the world supplies rather than one the plot needs. **"The story does not visit them yet" is not a reason.**
4. **★ Consent and capacity are decision-specific and time-specific.** A competent refusal on day three is not a status conferred for the rest of the episode. **When the facts change, the assessment renews and so does every duty resting on it.** If a person visibly worsens at lead 12, that is new information to everyone who can act, and the sweep runs again from there. A refusal the story keeps leaning on without re-earning it has stopped being a right the world supplies.
5. **★ Check causal order against your own dates.** If A is said to cause B, A's date must precede B's. **This sounds too obvious to write down. The first spine run under this sweep got it backwards,** with the antagonist's warning on day four closing a route the victim had already closed on day three.

Division of labour, by demonstrated performance

Not by preference. By what the record shows.

- **ChatGPT: audit and structure.** Its REWORK of `No Fixed Address v2.1` found six real defects including two I had made three times. Its seed ledger conformed where mine failed outright.
- **Fable: prose.** Wrote the best lines any of us has produced for this project. Needs an explicit brief for structure; give it a repaired spine and it will outwrite the rest of us.
- **Claude: spine, canon, integration, code.** Holds continuity across documents, catches contradictions, applies rulings, and does the engineering.
- **Stephen and Trish: the two questions only humans can answer.** Does it grab you, and do you believe it. Never conformance. Never volume.

PRE-FLIGHT: the thirteen objections

Every one is quoted from an objection actually raised by Stephen or Trish. This is empirical, not theoretical. Applied to a draft or a spine, never as a generation rule.

Answer each in one written sentence. An unanswerable item is a kill, not a note.

Believability

1. **Why was the decisive material not found at the time?** The answer may not be that anyone was lazy. *"Competent police would have done as a simple matter of process."* Legitimate: it did not exist then · the witness was a child · a promise released by a death · another jurisdiction · lawful destruction.
2. **Does the antagonist act in his own interest at every point?** *"Why would Dean call the workplace and leave his name?"*
3. **Would a real person actually do the small thing this needs them to do?** *"A mother would never put an admin form on a wall."*
4. **Is the motive one an ordinary reader accepts without argument?** *"Unbelievable motive for murder."*
5. **Is any discovery a coincidence wearing a coat?** *"I walked Rivermouth and found a landlord???*
Really?"

Comprehension

6. **Does a reader with no context understand every word?** *"What is a chandlery?"* No trade, legal or procedural vocabulary anywhere near the hook.
7. **Can a reader hold the cast in their head on one pass?** *"Too many characters to try and get my head around."*
8. **Is it clear what the crime is, immediately?** *"I don't understand what crime she is serving 11 years for."*

Knowledge and logic

9. **Can Ally source everything she says she knows?** *"How does she know about the mother? Did they meet?"* Every fact traces to a person, a document or a search, with a route.
10. **Does anything break its own rules?** *"Why would Ally say 'Wait. Where are you.' to a recording?"*
11. ★ **If the premise claims a route is the ONLY one, read your own access ledger back and find the second route.** *"Naming Kerry is therefore the safest official route, not the only lawful route, which makes the declared engine decorative."* **ADDED 2026-08-27, and it killed No Fixed Address at the spine.** The contradiction was on the same page: one section listed a person as reachable and willing to speak, the next section declared his route closed. **A premise whose engine is "this is the only way" is making a claim about everything it did not write down, and the cheapest place to falsify it is the ledger you just wrote.** If the wall has to be engineered back in, the engine was decorative and the premise is dead.

Stakes

12. **Is the stake a body, freedom, family, identity or present danger?** Never a licence, benefit, election, membership or apprenticeship. *"Unions are boring."*
13. **Does anything decisive rest on a small number?** *"Tills running nine minutes fast would just be synchronised."* Corroboration can make a correction credible; it cannot make arithmetic interesting.

One question that is not on the list, and should be asked anyway

Is the antagonist a person or a mechanism? The strongest predictor of true-crime consumption is interest in dangerous minds, not in bodily harm. Forty-five of our premises had antagonists who were situations. The audience wants access to a person.

What this replaces

Retire from generation: the twenty-four-seed fields, preallocated influence pairs, per-family distribution matrices, cross-audit instruments, and the twenty-field ledger schema. All were competently built. All were pointed at the wrong end of the process, and the field they produced was rejected wholesale by the only two readers who saw it.

Keep, unchanged: Kill Gate Zero and the least-cost test, the access map, the consequence-and-agency gate, the clock convention, evidential language, the canon roster and the bible. These are believability controls and they work. They belong in step 2 and step 3, applied to a draft.

The distinction that matters: a rule applied before writing shapes what gets written and produces conformance. The same rule applied after writing catches defects and costs the prose nothing.

What this cannot do

It cannot make a premise interesting. Nothing on this page would have rescued the union ballot or the till clock, and nothing here would have invented the mother still paying her daughter's phone bill.

What it does is stop us spending human attention on work that was never going to survive, so that the ten minutes we get from a reader lands on something that might.

The bottleneck is judgment, and judgment lives in two people. Everything above exists to feed them better and more often.